

Villain-Driven Irrationality: How Hating a Target Weakens Evidence, Justice, and Dissent

Phil Stilwell
Independent Scholar
ORCID: 0009-0009-3725-8682

Conceptual preview

VILLAIN-DRIVEN IRRATIONALITY

How perceived villainy erodes evidential standards

1 CORE MECHANISM



2 TYPICAL TACTICS



3 SOCIAL AMPLIFIERS



4 THE 'SECOND VILLAIN'



5 CONSEQUENCES AND SAFEGUARDS



“ Standards that work only in the absence of villains are not standards. ”

Abstract

This paper examines a common failure of judgment. Once a person or group has been cast as a villain, ordinary rules of evidence and fairness begin to look optional. I call this the *villain exception*. Its temptation is strongest when the target has done real wrong, because one proven offense can then become a blank check for believing the next rumor. Drawing on research about motivated reasoning, moral judgment, conformity, group identity, dehumanization, social cascades, and online outrage, the paper maps the tactics that follow: turning allegations into facts, treating association as participation, piling unlike claims together, moving a claim when challenged, reversing the burden of proof, and punishing people who ask for evidence. It also examines *connotative scattershot*, the use of many ugly labels in the hope that their combined stain will substitute for proving any one of them.

Two social mechanisms help the pattern persist. An *irrationality ratchet* forms when repeated shortcuts provide both apparent confirmation and social permission, while private doubts remain hidden. *Vice transfer* occurs when a person who restores an evidential distinction is treated as helping the villain, then as sharing the villain's allegiance and vices. The paper defines rational discipline in modest terms: match evidence to the exact claim, state confidence honestly, remain open to correction, and keep standards stable across targets. The villain exception damages public discussion, weakens due process, hides genuine wrongdoing inside clouds of exaggeration, and makes democratic safeguards depend on who is accused. Practical responses include separating claims, tracing source dependence, protecting good-faith dissent, distinguishing precaution from verdict, and applying the same rule after changing political jerseys. Rationality is tested most severely when accuracy happens to help someone we have good reason to despise.

Keywords

villain exception; motivated reasoning; social validation; informational cascades; pluralistic ignorance; vice transfer; moral outrage; partisan bias; due process; evidential standards; moral typecasting; belief in pure evil; affective polarization; dehumanization; connotative scattershot; public discourse

1. Introduction

Most people endorse evidential caution in the abstract. Serious accusations should be supported. Distinct claims should not be merged. Allegations, suspicions, corroborated findings, civil judgments, and criminal convictions should not be treated as interchangeable. A person's association with a wrongdoer is not automatically participation in the wrongdoing. A demand for evidence is not a defense of the alleged act. These propositions sound elementary when no emotionally protected conclusion is at stake.

The test arrives when the accused is already hated.

Suppose a public figure has lied, abused power, behaved cruelly, or been credibly implicated in serious misconduct. The rational response is not to forget those facts. Prior conduct may alter the probability of further misconduct, provide context, justify investigation, and reduce confidence in the person's denials. Yet another inference often appears. Because the person is bad, a further accusation is treated as true without evidence adequate to that accusation. Because the person is dangerous, distinctions among kinds of wrongdoing are treated as evasions. Because the person's allies have often misled the public, every denial is taken as confirmation. Because the stakes are high, the safeguards designed for high-stakes judgment are suspended.

Prior wrongdoing is like a smoke alarm. It gives us a reason to look quickly and carefully. It is not security-camera footage of the next event. If an official has lied ten times, an eleventh claim deserves extra checking. The first ten lies do not tell us, by themselves, whether the eleventh claim is false, half true, or fully true.

This paper examines that inversion. The existence of a perceived villain does not merely supply a conclusion. It can alter the rules by which conclusions are reached. Standards that were represented as universal become conditional. The accused must disprove what the accuser has not established. Association becomes complicity; silence becomes confession; denial becomes further evidence; a request for a source becomes proof of moral corruption. The worse the target is believed to be, the less accurately the target must be described.

I call this pattern the *villain exception*: the tacit permission to weaken evidential and procedural standards when a target has been assigned sufficiently negative moral status. Its psychological signature is a correlative decline in rational discipline. As perceived villainy rises, the threshold for accepting condemnatory claims falls, the threshold for accepting exculpatory information rises, and the social cost of requesting precision increases.

That last cost introduces a second target into the drama. The first target is the villain. The second is the person who interrupts the accelerated judgment and asks which proposition the evidence establishes. This person may fully condemn the villain's documented conduct and still be treated as a partisan of the villain. The progression is remarkably economical: to question an allegation is to impede condemnation; to impede condemnation is to aid the accused; to aid the accused is to share the accused's allegiance; and to share the allegiance is to possess the accused's vices. The intervening rational mind becomes morally assimilated to the object it is trying to reason about.

The vice assigned through this process need not have been established in the original target. If the villain is merely rumored to be treacherous, predatory, hateful, or fanatical, the person asking whether the rumor is true can be accused of the same trait. An unproved vice is copied from an accused target to an evidential critic without acquiring proof at either location. The result is a peculiar moral multiplication: uncertainty about one person's conduct produces certainty about two people's character.

A schoolyard version makes the move easier to see. One child says, "Sam stole the phone." Another asks, "Did anyone see Sam take it?" The crowd answers, "Why are you protecting a thief?" The second child has not defended stealing or even said that Sam is innocent. A request for a witness has been changed into evidence of bad character.

The exception is also social before it becomes stable. One person's overstatement can be dismissed as a lapse. A hundred parallel overstatements look like a norm. Repetition supplies a descriptive message about what members do, approval supplies an injunctive message about what members ought to do, and the combination changes the cost of independent judgment. Later speakers do not merely copy a conclusion. They learn the acceptable degree of certainty, the approved insults, the permitted shortcuts, and the penalty for restraint. Because each public performance becomes evidence to the next observer, irrationality can acquire momentum even when many participants privately retain doubts.

The claim is not that every person described as a villain is misunderstood. Some leaders, movements, institutions, and private actors are dangerous. Some accusations that circulate in crude or partisan forms are nevertheless true. Nor does the argument imply that all political factions produce equal quantities of falsehood in every period, on every subject, or through institutions of equal power. Political symmetry at the level of human susceptibility does not entail historical or empirical equivalence in every case.

The narrower claim is about the rules of judgment: no person's badness can establish a fact that the evidence does not establish. A documented offense may make a related offense more likely; it does not prove every nearby charge. Moral urgency can justify quick investigation and temporary precautions. It does not erase the difference between what is possible, likely, well supported, and formally decided. Asking for support

for a particular claim is not a favor granted to admirable defendants. It is part of caring whether accusations are true.

The argument draws its taxonomy from recurring patterns in ordinary online disputes, but those disputes are not treated as scientific data and no participant is analyzed. The examples below are deliberately generic and politically reversible. The psychological claims instead draw on experimental and observational research concerning motivated reasoning, moralization, partisan evaluation, dehumanization, outrage, and social media. The paper is therefore a conceptual synthesis and normative analysis, not a new causal study.

Its contribution is to join four levels that are often discussed separately: the lowered evidential threshold applied to a hated target, the tactics that protect weak claims from correction, the group processes that make those tactics feel normal, and the transfer of the target's vices to anyone who interrupts the pattern. The resulting framework generates both practical safeguards and testable predictions. It does not assume that all mechanisms operate in every dispute or that every harsh accusation is a product of them.

The structure is as follows. Section 2 defines the villain exception, distinguishes legitimate character evidence from a general evidential waiver, and models the transfer of vice to a rational intervener. Section 3 develops the psychological and social architecture of the exception. Sections 4 and 5 identify its evidential and coercive tactics. Section 6 explains how those tactics acquire sociological inertia through social validation, threshold effects, informational cascades, pluralistic ignorance, and escalating norm performance. Section 7 analyzes connotative scattershot, the practice of throwing maximally negative labels until one adheres. Sections 8 through 10 examine the effects on justice, public discourse, and society. Section 11 proposes disciplines for rational judgment in the presence of a villain. Section 12 answers objections, and Section 13 outlines an empirical research program. Three appendices provide a practical accusation audit, a script-to-safeguard crosswalk, and a reusable research vignette and coding template.

2. The Villain Exception

2.1 What rational discipline means here

The paper uses *rational discipline* in a modest, procedural sense. A person reasoning rationally should identify the claim being judged, match reasons to that claim, distinguish background suspicion from case-specific support, state confidence at the level the evidence warrants, and remain able to say what would change the judgment. The same basic rule should survive a change in the target's identity.

This does not require emotional neutrality, moral hesitation, formal proof in every conversation, or equal suspicion toward unequal facts. Anger can reveal that something important is at stake. A terrible record can raise a reasonable prior suspicion. Urgent risk can justify protective action before a final finding. Rational discipline governs the next step: what the emotion, history, or risk permits us to claim as true.

The word *villain* also names a role in a judgment, not a claim that the target is innocent or merely misunderstood. A target may fully deserve severe condemnation. The question is whether that condemnation becomes permission to loosen the connection between a further claim and its support.

2.2 Bad character is not a universal premise

A person's established conduct can be evidentially relevant. If an official has repeatedly falsified records, that history can rationally affect an assessment of a new record. If an organization has concealed abuse, that record may justify closer scrutiny of its present disclosures. If a speaker has lied about a topic, confidence in an unsupported statement on that topic should decrease. Rational analysis does not require pretending that every claim arrives without a history.

But relevance has limits. Character evidence usually changes the plausibility of a claim; it does not supply the missing elements of the claim. Evidence that a leader is corrupt is not automatically evidence of a specified financial transaction. Evidence that a movement tolerates extremists is not automatically evidence that every member endorses the extremists' acts. Evidence of adult sexual misconduct is not, without an additional bridge, evidence of misconduct involving children. Evidence that an administration resisted disclosure is not by itself proof of what every withheld document contains.

The villain exception erases these limits. It converts a justified global judgment - "this person is dangerous, dishonest, abusive, or corrupt" - into an unrestricted local premise: "therefore this further accusation is true." The target's established wrongs become an evidential warehouse from which an accuser may withdraw support for any conclusion without identifying the relevant item.

Consider a reckless driver. A long record of speeding is a good reason to inspect that driver's role in a new crash. It is not proof that the driver ran this red light. Investigators still need the light sequence, witness accounts, vehicle data, or video. Bad history can point the flashlight. It cannot stand in for what the flashlight is meant to find.

The problem is not harsh criticism. Severe facts warrant severe language. The problem is using moral status as a missing bridge. A conclusion is not proved merely because it would fit a person's character. Fit can guide the search. It cannot finish it.

2.3 A pseudo-formal model

Let c be a condemnatory claim about target t . Let $Q(c)$ represent the observed quality of the claim-specific support. Let $T(c)$ represent the level of support required for the confidence being asserted in that setting. The threshold may reasonably reflect the claim's specificity, the consequences of error, prior probability, and whether the speaker is expressing suspicion, publishing a finding, or imposing a penalty. Acceptance is warranted when $Q(c) \geq T(c)$. Let $V(t)$ represent the target's perceived villainy.

The villain exception changes the perceived threshold:

$$\hat{T}_{condemn}(c, t) = T(c) - \lambda V(t)$$

where λ is the degree to which perceived villainy discounts the evidence required for a condemnatory conclusion. For potentially exculpatory claims, the direction reverses:

$$\hat{T}_{exculpate}(c, t) = T(c) + \mu V(t)$$

where μ represents the additional evidence demanded before information favorable to the villain is accepted. The *villain discount* is therefore:

$$VD(c, t) = T(c) - \hat{T}_{condemn}(c, t)$$

This is not a measured law of the mind. It is a simple picture of a thumb on the scale. The problem is not that $V(t)$ can affect a reasonable prior probability. The problem is that it discounts the amount of case-specific support treated as necessary. One weak report may be enough for condemnation, while several corrections are still not enough for revision. The same person who checks the source, setting, and legal status of an allegation against an ally may treat an allegation against an enemy as proof of itself.

The paper's central claim about declining rational discipline can be represented equally simply. Let $R(t)$ denote adherence to stable rational standards when evaluating claims about t . The hypothesis is:

$$R(t) = R_0 - \alpha V(t) \text{ when the villain exception is active.}$$

Again, the equation is a sketch, not a measured result. Its point is that villainy should be something we judge, not acid that dissolves the rules we use to judge.

2.4 The two-error discipline

Every accusation creates at least two dangers. A false negative fails to identify a real wrongdoer. A false positive condemns a person for something not adequately established. Public argument often moralizes only the first error. Caution is portrayed as indifference to victims, while overstatement is excused as excessive concern.

That asymmetry is unstable. False accusations can harm the accused, but they also harm genuine victims. They dilute reliable testimony, give wrongdoers examples with which to discredit sound allegations, consume investigative attention, and train audiences to treat correction as partisan maneuver. Conversely, fear of false accusation can be abused to silence complainants. Rational standards must therefore manage both errors rather than sanctify either credulity or disbelief.

A high cost of false rejection may justify precautionary action before final adjudication. An institution may separate parties, preserve records, suspend access, or initiate an independent investigation. Precaution, however, should be described as precaution. The need to reduce risk does not transform an allegation into a finding. Action under uncertainty and belief under uncertainty are related but distinct.

A fire alarm offers a useful analogy. We evacuate before proving that the building is burning because the cost of waiting may be terrible. But the evacuation itself is not proof of a fire. We still inspect the building. In the same way, a school or workplace may act to protect people while an allegation is investigated. The protective step should not later be cited as proof that the allegation was already settled.

2.5 The villain can be guilty

The villain exception is easiest to see in a fabricated accusation against a blameless person, but that is not its most important case. Its strongest test is a target who has done something genuinely terrible. Rationality requires neither sympathy nor neutrality about established wrongdoing. It requires accuracy about what the wrongdoing was.

If a target has committed offense A, one may condemn A without hesitation. If A raises the probability of B, one may state that B deserves investigation. If evidence supports B, confidence should rise in proportion to that evidence. What one may not do is announce B as fact merely because A makes B emotionally satisfying.

This restraint is sometimes described as "defending" the villain. The description assumes that evidential standards are personal benefits rather than public rules. But rules of evidence do not exist chiefly to protect the reputations of good people. They exist because humans are unreliable judges under pressure, because power can manufacture accusations, and because even a contemptible person can be falsely accused of a particular act.

A referee does not endorse a dirty team by calling the ball out of bounds correctly. The rule matters most when everyone in the stadium wants the call to go the other way. In the same fashion, refusing one false charge against a guilty person does not erase the true charges.

2.6 The intervening rational mind and vice transfer

The villain exception is maintained not only by conclusions about the target but also by conclusions about anyone who interrupts those conclusions. Let i be an intervener who questions claim c about villainized target t . A defensible response evaluates the intervention itself: Is the objection relevant? Is its evidential standard attainable? Is it applied consistently? Does it acknowledge contrary evidence? The villain exception instead uses a sequence of identity transformations:

$$questions(i, c) \rightarrow aids(i, t) \rightarrow aligns(i, t) \rightarrow inheritsVice(i, t)$$

The arrows do not show valid steps. They show a social shortcut. A question may delay a verdict, and delay may happen to benefit a guilty target. But helping someone in one limited way does not prove loyalty to that person, much less a shared character. A defense attorney can insist on lawful procedure without approving a client's conduct. An editor can reject an unsupported accusation without admiring its subject. A scientist can correct a false claim about a dangerous product without declaring the product safe. We should judge the work being done, not simply notice who benefits from it.

Call the final move *vice transfer*: giving the intervener a vice assigned to the villain simply because the intervener restored a distinction that helps the villain in this dispute. It has two forms. In *established-vice transfer*, a real defect of the target is copied onto the critic without evidence about the critic. In *imagined-vice transfer*, even the target's defect has not been proved. The rumor is copied anyway. It is like using a blurry photograph to make a second blurry photocopy, then pointing to the two images as mutual confirmation.

The mechanism is strongest when the intervener shares a group identity with the accusers. An outsider can be dismissed as an enemy. An insider who insists on claim-specific evidence threatens to show that membership does not require credulity. The intervention therefore carries information beyond its words: another norm is possible. Silencing or contaminating the intervener prevents that information from becoming common knowledge.

There are genuine cases of complicity. A supposed neutral may destroy records, intimidate witnesses, coordinate a false account, demand evidence no one could ever provide, or use procedure only to run out the clock. Such conduct should be named and supported. Vice transfer begins when the verdict comes first - when asking for evidence is itself treated as evidence of corruption. The key question is ordinary: What did this intervener actually do? It is not enough to answer, "Something that happened to help a bad person."

3. The Psychological Architecture of the Villain

3.1 Directional reasoning and unequal thresholds

Kunda's (1990) account of motivated reasoning distinguishes an accuracy motive from a directional motive. When accuracy governs, people try to use procedures they regard as appropriate to reach a correct conclusion. When a preferred conclusion governs, they recruit memories, beliefs, and strategies that can support that destination, subject to the constraint that the justification must appear plausible.

The villain exception supplies a powerful directional goal. Condemnation protects the perceiver's identity, coalition, and moral self-conception. Acquittal, uncertainty, or even terminological correction can feel like a concession to evil. The task silently changes from "What does this evidence establish?" to "How can this material support the verdict that a bad person deserves?"

Ditto and Lopez (1992) demonstrated a closely related asymmetry in decision thresholds. Participants scrutinized unwelcome conclusions more critically and required more information before accepting them than

they required for preferred conclusions. Taber and Lodge (2006) found motivated skepticism in political evaluation: participants counterargued attitude-incongruent material, accepted congenial arguments more readily, and selected confirmatory information when given a choice. The villain discount is a moralized version of this differential processing.

The mechanism need not involve deliberate lying. A person can sincerely experience a weak claim as decisive because scrutiny has been allocated asymmetrically. The condemnatory item receives a quick plausibility check. The correction receives adversarial cross-examination. The individual then reports the products of unequal procedures as though they were outputs of one standard.

It is as if two lawyers live in the same head. The claim we want to believe gets a friendly lawyer who asks, “Can I make this work?” The claim we dislike gets a hostile lawyer who asks, “Can I find any flaw?” Both lawyers may use real facts. The bias lies in which lawyer is assigned to which claim.

3.2 Moral construal, extremity, and universality

Representing conduct as a moral matter can accelerate and intensify evaluation. Van Bavel, Packer, Haas, and Cunningham (2012) found that moral construal produced faster, more extreme, and more universal evaluations of the same actions than pragmatic or hedonic construal. Moral cognition is not therefore irrational. Some actions should be evaluated morally, and moral clarity can support courage. The danger is that moral construal compresses degrees, contingencies, and alternative explanations just when a claim requires careful discrimination.

Once a target is cast as a villain, the question “Did this person commit this particular act?” becomes fused with “Are you against evil?” The first is empirical and specific. The second is expressive and identity-laden. A speaker who answers the first with uncertainty can be heard as answering the second with indifference.

This fusion helps explain why evidential precision is treated as moral evasion. Distinguishing one offense from another sounds like weakening condemnation, even when the first offense is explicitly accepted. The morally saturated audience hears gradation as exoneration because the desired message is total opposition, not accurate classification.

Moral framing can work like turning the volume knob all the way up. The louder setting may be fitting when the danger is real. But at full volume, small differences become hard to hear. “I am not sure this detail is true” is mistaken for “I do not care whether anyone was harmed.”

3.3 Pure evil and the disappearance of situation

Webster and Saucier (2013) operationalized individual differences in belief in pure evil. Across five studies, stronger belief in pure evil was associated with more aggressive intergroup orientations and less support for rehabilitative or prosocial responses. Their work is especially relevant because the pure-evil representation treats harmful character as stable, intentional, and largely beyond situational explanation.

A villain constructed in this form is not simply a person who has done wrong. The villain is wrongdoing embodied. Context becomes excuse; explanation becomes sympathy; patience becomes collaboration. Because the villain’s essence is presumed to generate harm, almost any further harm can be attributed without a separate causal account.

This essentialism also reduces the value of correction. If the target is corrupt in essence, a false accusation can be dismissed as true “in spirit.” The accuser may concede that a specific detail was wrong while retaining the moral satisfaction of the original claim. Accuracy becomes secondary to expressive fit.

The villain then becomes a stencil placed over every new event. Whatever happened underneath is traced into the same outline. A lost document, a clumsy statement, and a deliberate cover-up may all come out looking identical because the stencil, not the event, supplies the shape.

3.4 Moral typecasting and lost victimhood

Gray and Wegner (2009) found an inverse relation between perceived moral agency and perceived moral patiency. People strongly represented as moral agents - including evil-doers - were seen as less capable of being wronged. This phenomenon, called moral typecasting, illuminates the justice problem at the center of the villain exception.

A person can be both an agent of wrongdoing and a patient of further wrongdoing. A corrupt official can be defamed. An abusive person can receive an unfair proceeding. A violent movement can be blamed for an act it did not commit. These propositions are conceptually simple, but moral typecasting makes them psychologically difficult. Once a target occupies the villain role, harm done to the target no longer registers fully as harm.

The result is a dangerous inference: because the person deserves punishment for something, it matters less whether the present accusation is true. Desert becomes fungible. An unsupported allegation is excused on the ground that the target has earned condemnation elsewhere.

A reckless driver can still be rear-ended while stopped at a light. The driver's bad record does not make the new collision imaginary, and it does not tell us who caused it. In the same way, a wrongdoer does not lose the ability to be wronged.

3.5 Moral disengagement and authorized harshness

Bandura's (1999) theory of moral disengagement describes mechanisms by which people suspend self-sanctions against harmful conduct. These include moral justification, sanitizing language, advantageous comparison, diffusion of responsibility, minimization of harm, attribution of blame, and dehumanization. The villain exception recruits several of these mechanisms at once.

Smearing the target is redescribed as protecting victims. Professional punishment is redescribed as accountability. A false or exaggerated accusation is compared with the target's worse documented conduct and therefore made to seem negligible. Responsibility for accuracy is diffused across the crowd: everyone has heard the claim, so no single repeater feels obliged to verify it. The target is blamed for producing the conditions under which people lie about the target.

The reasoning resembles a group passing around a damaged tool. Each person says, "I only used what everyone else was using." Because responsibility is spread across many hands, no one feels responsible for the bad result. Crowds can divide responsibility without dividing the harm.

Dehumanization is particularly consequential. Kteily, Bruneau, Waytz, and Cotterill (2015) found that blatant dehumanization predicted support for aggressive actions, including retaliatory violence and torture, beyond ordinary prejudice measures. Villain language need not become openly dehumanizing in every case, but it moves in that direction when a target is represented as contaminating, monstrous, verminous, or outside the class of persons to whom procedural concern is owed.

3.6 Coalition, affect, and symmetry

The villain is often the personification of an outgroup. Iyengar and Westwood (2015) found that partisan cues shaped judgments and discriminatory behavior well beyond policy disagreement. Affective polarization

turns political identity into social aversion. Under these conditions, an accusation against an opponent is not merely information about an individual. It is a resource in intergroup competition.

The susceptibility is not owned by one side. Ditto and colleagues (2019) meta-analyzed 51 experimental tests involving more than 18,000 participants. They found robust partisan bias and no meaningful mean difference between liberals and conservatives in the tendency to evaluate otherwise identical information more favorably when it supported their political commitments. This does not show that every faction is equally inaccurate on every issue. It shows that directional evaluation is a broadly distributed human capacity.

Political balance in this paper is therefore not a ritual of splitting blame. It is an experimental constraint. Any proposed safeguard should survive a reversal of names, parties, causes, and protected identities. If a standard is compelling only when one's opponent is accused, it is not yet a standard.

This is the jersey-switch test. Keep the play the same and change the team colors. If a foul becomes fair play when our side commits it, we are judging jerseys rather than conduct.

3.7 Normative influence, informational influence, and social permission

Individual motivation cannot explain the durability of the villain exception by itself. People judge in view of other judges. Deutsch and Gerard (1955) distinguished normative social influence - influence produced by the desire to meet others' expectations - from informational social influence, in which others' judgments are treated as evidence about reality. In villain-centered argument the two are easily fused. A person sees that condemnatory certainty is rewarded and learns both that certainty is socially approved and that so many confident people must know something.

Cialdini, Reno, and Kallgren (1990) further distinguished descriptive norms, which communicate what people do, from injunctive norms, which communicate what people approve. A stream of accusatory posts, slogans, denunciations, or professional declarations can carry both messages. Their frequency says that members repeat the claim; visible praise and the punishment of qualification say that members ought to repeat it. The observer who conforms then increases the apparent frequency and helps enforce the approval rule. Social validation is therefore recursive.

A line outside a restaurant offers a harmless example. The line tells a passerby that many people chose the place, so the food may be good. It also makes joining feel normal. But if everyone joined because they saw the same line, the crowd does not contain one hundred independent restaurant reviews. Political crowds can grow in the same way, with much higher stakes.

This recursion can change not only whether a proposition is voiced but how irrationality is performed. A community learns that bare assertion is tolerated, that a certain family of labels is rewarded, that sources are optional for allies, and that procedural distinctions invite suspicion. Participants calibrate to the room. A modest allegation can signal insufficient commitment after stronger accusations become common, so a member matches or exceeds the prevailing intensity to make allegiance legible.

Moscovici and Zavalloni (1969) found that group discussion could move attitudes toward a more extreme position, with the shift persisting in later individual judgments. The present argument does not claim that every discussion polarizes or that extremity is always false. It identifies a special danger when the direction of group movement is also treated as evidence of its correctness. More severe rhetoric then performs three functions at once: it states a view, signals a loyal identity, and alters the benchmark against which the next speaker's loyalty will be judged.

3.8 Moral rebels, stigma by association, and the second villain

Why attack a person who voices a sound standard? Because the correction can function as an implicit reproach. Monin, Sawyer, and Marquez (2008) found that people who had complied with objectionable instructions could resent a principled refuser whose conduct threatened their positive self-image. A rational intervener in a villain narrative can pose an analogous threat. The intervener's restraint suggests that the group could have opposed the villain without lowering its standards. What sounds like pedantry may therefore be experienced as an accusation against everyone who already assented.

Imagine a group project in which everyone copies an answer except one student. The student need not lecture anyone. Merely refusing to copy can make the others feel judged. Likewise, a person who calmly asks for a source may remind the group that its certainty was a choice, not a duty.

Pryor, Reeder, and Monroe (2012) demonstrated stigma by association: companions of stigmatized people can themselves be devalued, through both deliberative and spontaneous processes. The villain exception extends the category of "companion" beyond actual relationship. A person can be made socially adjacent to the villain by defending a procedure, correcting a quotation, challenging a diagnosis, or refusing a label. The conversational association is then mistaken for a moral one.

The group-protective motive can be especially strong toward an internal dissenter. Marques, Yzerbyt, and Leyens (1988) found more extreme judgments of desirable and undesirable ingroup members than of comparable outgroup members, a pattern they called the black sheep effect. The result does not by itself establish how contemporary political groups treat evidential critics. It supplies a relevant mechanism: an ingroup deviant can threaten the group's positive identity precisely because the deviant is recognizably one of "us." The rational intervener is dangerous not because one more opponent has appeared, but because a loyal member has demonstrated that loyalty and credulity can be separated.

These processes create the *second villain*. The first villain organizes the original condemnation. The second organizes discipline inside the condemning group. Traits assigned to the first - cruelty, corruption, prejudice, treachery, predation, fanaticism - are copied onto the second, even when the second accepts the accurate charges and disputes only the unsupported ones. Once one critic has been publicly punished, the crowd does not need to threaten every doubter. The example works like a warning sign at the edge of the road.

4. Tactics of Evidential Degradation

The villain exception has a familiar toolbox. Its tactics often work together, so several weak moves can create one strong feeling.

4.1 The unlocated evidential warehouse

An accuser is asked for support for a specific claim and responds by gesturing toward a large scandal, archive, investigation, movement, or body of reporting. The scandal may be real and the archive may contain important evidence. But the requested link is not identified.

This maneuver transfers the labor of proof to the challenger. The challenger is invited to search a vast body of material and assume that the relevant support must be present somewhere. If the challenger cannot refute every item in the warehouse, the original claim is treated as intact.

Rational inquiry works in the opposite direction. The claimant identifies the source, the proposition it supports, and the inferential path from that proposition to the conclusion. A large quantity of adjacent material is not a substitute for a relevant item.

Imagine asking a librarian for the page that supports a quotation and being told, “It is somewhere in this building.” A library may contain a million true statements, but its size does not locate this one. The person making the claim must at least point to the right shelf.

4.2 Allegation-to-fact conversion

The existence of an allegation is evidence that an allegation was made. It is not by itself evidence that the event occurred. This distinction can sound pedantic only because ordinary discourse frequently erases the reporting verb. “A witness alleged X” becomes “documents reveal X,” and then “we know X.”

The conversion often proceeds by repetition rather than discovery. A claim is copied across accounts, summarized by commentators, mentioned by an official, and embedded in a compilation. Each retelling appears to be another source even when all versions trace back to one unverified statement.

Fazio, Brashier, Payne, and Marsh (2015) found that repetition increased perceived truth even when participants possessed relevant knowledge. Fluency can therefore imitate verification. A repeatedly encountered allegation feels familiar, and familiarity is misread as evidential depth.

Social repetition adds another illusion: apparent independence. Ten people may seem to supply ten judgments even when nine inferred truth from the confidence of the first or copied one another without checking the source. Headcount has evidential value only to the extent that the judgments reflect sufficiently independent access or analysis. When the same rumor is socially learned down a chain, the crowd is not ten witnesses. It is one uncertain signal with nine echoes.

Ten photocopies of one receipt are not ten purchases. In the same way, ten accounts that all trace back to one claim are still one source line, however crowded the screen looks.

4.3 Possibility, probability, and certainty collapse

Villain reasoning moves easily among modal levels. A target *could* have committed an act. Existing conduct makes it *plausible*. Several suspicious facts make it *likely*. The claim is then stated categorically. When challenged, the speaker retreats to the earlier and weaker proposition that the act remains possible.

Each proposition requires a different defense:

- possibility requires compatibility with known facts;
- plausibility requires a credible mechanism or fit with background evidence;
- probability requires comparative support over alternatives;
- categorical assertion requires a level of confidence appropriate to the context and stakes.

Moving among these levels is not a minor wording issue. It makes a claim difficult to test because criticism is directed at the strong version while defense is offered for the weak one.

The move is like climbing a ladder from “could” to “did” and then jumping back to the bottom whenever someone checks the top rung. The speaker gets the force of certainty and the safety of possibility.

4.4 The evidential heap

An evidential heap combines heterogeneous items: a documented offense, an association, a rumor, a disputed document, a suspicious decision, an inaccurate legal claim, a personality judgment, and an argument from silence. The pile is then assigned a collective name such as “overwhelming evidence.”

Aggregation can be rational when independent pieces support the same proposition. The heap is irrational when the pieces support different propositions or when their dependence is concealed. Evidence that a person lied about one matter, knew a wrongdoer, resisted disclosure, and behaved abusively elsewhere may jointly raise suspicion. It does not automatically establish a specified act unless the items bear on that act.

The heap also creates a defensive burden trap. A critic who corrects one item is accused of ignoring the rest. A critic who addresses every item produces a long answer that can be mocked as obsessive or evasive. Quantity becomes both the substitute for relevance and the means of punishing anyone who restores distinctions.

A grocery bag full of unrelated objects does not become a meal because it is heavy. In the same way, a large pile of facts is not automatically a case. The items must connect to the same claim.

4.5 Claim drift and the strategic retreat

When a narrow accusation is challenged, it broadens. A claim that the target personally committed act A becomes a claim that the target enabled a culture in which A occurred, protected someone who committed A, resisted an inquiry into A, or simply belongs to a class associated with A. These broader claims may deserve investigation, but they are not the original claim.

Once the broader proposition is defended, the speaker returns to the narrow wording as though nothing changed. This is a practical form of motte-and-bailey reasoning. The emotionally powerful claim attracts attention; the easier, weaker claim provides shelter.

It resembles a shell game. The dramatic claim draws the eye, but when challenged the argument slides under a safer shell. Once attention moves away, the dramatic claim appears again.

4.6 Guilt by association and moral contagion

Association can be evidence when the nature, duration, context, and purpose of the association are specified. A financial partnership, command relationship, conspiracy, or documented exchange can establish relevant connections. But mere proximity is not participation.

Villain reasoning treats moral status as contagious. Contact becomes knowledge; knowledge becomes approval; approval becomes assistance; assistance becomes commission. The steps are compressed because each intermediate proposition feels consistent with the target's character.

The same contagion spreads to interlocutors. A person who requests evidence is said to defend the target. A person who defends the target's procedural rights is said to endorse the alleged conduct. A person who objects to a label is placed inside the labeled category. Association replaces analysis at every level.

Vice transfer is the character-level completion of this chain. The intervener does not merely become associated with the target; the target's moral inventory is copied onto the intervener. This is most conspicuous when a critic corrects one inflammatory label and is immediately called by that label. The accusation has become self-protecting: anyone who asks whether the category fits is offered as a second example of the category.

The logic treats vice like wet paint. Touch a rule that happens to protect the villain and the villain's color is supposed to rub off on you. Real responsibility does not work that way. It depends on what a person knew, intended, and did.

4.7 Burden reversal and the proof of innocence

The accuser may demand that the target or critic demonstrate that the allegation is false. This is sometimes legitimate when the target possesses uniquely accessible records or has a duty of disclosure. It is not a general replacement for the accuser's burden.

Negative claims can be difficult or impossible to establish, especially when the accusation is vague. A demand to disprove all possible versions of a rumor makes accusation costless. Rational allocation of burden begins with the party asserting a specific proposition, then shifts as evidence and access justify.

The classic absurd example is an invisible animal in the room. If one person claims it is there, everyone else cannot reasonably be required to prove that no invisible animal of any kind is hiding anywhere. The first burden belongs to the person making the claim.

4.8 Silence, litigation, and response as oracles

Every possible response can be made incriminating. Silence is guilt. Denial is panic. A detailed answer is overcompensation. A short answer is evasion. A lawsuit shows intimidation. Failure to sue shows fear of discovery. Cooperation shows image management. Resistance shows concealment.

The same trap often arrives in sharper language: *"Silence tells us everything," "That long denial is exactly what a guilty person would write,"* or *"If the charge were false, the target would have sued."* Each line pretends to read one response while quietly making every response point in the same direction.

This is not cumulative evidence. It is interpretive overfitting. A response has evidential value only relative to credible alternative expectations. If every response confirms the same conclusion, the accusation system is not learning from behavior.

Such a test is a two-headed coin: heads means guilty and tails also means guilty. It looks like a test because the coin is tossed, but the result was fixed before the toss.

4.9 Publicity as authentication

A public official, journalist, scholar, influencer, or large audience repeats the allegation. The fact of repetition is then treated as authentication. Expertise can matter, but an expert conclusion has evidential force only when the relevant expertise, access, method, and underlying reasons are identified.

An unnamed or merely prominent authority often functions as a prestige bridge across an evidential gap. Public confidence establishes confidence. It does not establish the event.

A microphone makes a voice louder; it does not make the sentence truer. Fame, office, and repetition can amplify a claim without adding another fact.

4.10 The genre exemption

When a meme, joke, slogan, or clipped video communicates a factual accusation, challenge is sometimes met with the claim that the genre is unserious. Yet the social effect depends on the audience receiving the implication as serious enough to condemn the target.

Humor permits compression and exaggeration, but it does not eliminate truth conditions from factual insinuation. One cannot rely on an audience to believe an allegation and then invoke entertainment when asked to support it. The genre exemption is a one-way valve: seriousness for persuasion, playfulness for accountability.

The move is a “believe me unless challenged” card. The meme is serious while it harms the target and suddenly only a joke when evidence is requested.

A typical retreat sounds like this: “*It was only a joke - but if the joke made you angry, maybe it landed too close to the truth.*” The speaker first refuses responsibility for the claim, then treats the objection as fresh proof of it.

5. Tactics of Social Closure

Weak evidence does not always survive through inference alone. It is often protected by making rational scrutiny socially costly.

5.1 The loyalty test

The basic loyalty test offers two positions: accept the accusation or reveal sympathy for the villain. It excludes the position that one can oppose the target, condemn established conduct, support investigation, and still reject an unsupported claim.

This false dilemma is powerful because it attaches epistemic caution to moral identity. The audience is not asked only what it believes. It is asked what kind of people its members are. Public assent becomes a badge of decency; qualification becomes a stain.

The loyalty test also separates private belief from public expression. A person may privately doubt a claim but repeat it to avoid suspicion. The resulting apparent consensus then pressures others to do the same. What looks like independent agreement may be coordinated fear.

The test builds a hallway with only two marked doors: “believe the charge” and “join the villain.” A third door - “oppose the villain and check this claim” - is painted over.

5.2 Complicity by scrutiny

A request for sources is recoded as assistance to the accused. A distinction among offenses is called minimization. A correction of legal or clinical terminology is described as concern for reputation over concern for victims. The substance of the scrutiny disappears, replaced by an allegation about its social function.

Scrutiny can, of course, be performed in bad faith. Endless demands, impossible standards, selective doubt, and harassment of complainants can obstruct truth. But bad-faith scrutiny must be shown through its pattern: standards no evidence could satisfy, refusal to acknowledge corroboration, selective burdens, or strategic delay. It cannot be inferred merely from the fact that scrutiny inconveniences a preferred conclusion.

A brake can be misused to stall a journey, but that does not make every brake a form of sabotage. Sometimes slowing down prevents a crash. We should inspect how the brake is used instead of condemning the brake because someone disliked the delay.

5.3 Vice transfer and the derivative villain

Complicity by scrutiny initially alleges a harmful function: the question helps the villain. Vice transfer turns that functional claim into a character claim. The questioner is said to possess whatever vice makes the villain hateful. A correction about a charge of treachery proves the critic disloyal; a request for support for a charge of prejudice proves the critic prejudiced; a distinction concerning alleged predation proves the critic friendly to predators.

The move depends on moral transitivity:

$$attributesVice(t, v) + apparentAid(i, t) \rightarrow attributesVice(i, v)$$

Neither premise necessarily holds, and the conclusion does not follow even if both do. A person can perform an act that benefits a wrongdoer for independent, principled, negligent, or self-interested reasons. More importantly, the relevant “assistance” may consist only in preventing a false description from being added to a true indictment.

The transfer becomes socially effective through compression. Instead of answering a claim about evidence, the crowd announces what sort of person would raise it. The intervener must now defend an entire identity while the original evidential question goes unanswered. Any persistence appears obsessive; any anger appears revelatory; withdrawal appears guilty. The second accusation system inherits the closed structure of the first.

Vice transfer often comes packaged as a sneer: *“Only someone who shares the vice would work this hard to correct the accusation.”* The sentence does not answer the correction. It makes the act of correcting into evidence of bad character.

The generic pattern is politically reversible. A person can reject a false allegation about a right-wing target and be assigned that target’s alleged bigotry. Another can correct a false allegation about a left-wing target and be assigned that target’s alleged extremism. A religious critic can be called impious for opposing a rumor about a religious leader; a secular critic can be called doctrinaire for opposing a rumor about a secular institution. The labels vary, but the function is constant: make epistemic intervention personally expensive.

The referee analogy is again useful. Overturning one bad call may help a team known for dirty play. That does not put the referee on the team’s roster. The call should be judged by the replay, not by the color of the jersey that benefits.

5.4 The closed accusation system

An accusation system is closed when every possible answer confirms the charge. Denial reveals defensiveness. Evidence against the claim reveals a cover-up. Neutrality reveals hidden allegiance. A demand for falsification criteria reveals moral coldness. The target’s guilt and the critic’s corruption become mutually supporting premises.

Closed systems imitate strength because they are difficult to refute. Their real defect is the opposite: they cannot distinguish worlds in which the accusation is true from worlds in which it is false. A claim that predicts every observation predicts nothing.

A practical test is simple: What possible response, document, or event would reduce confidence in the allegation? If the answer is “none,” the belief is functioning as an identity commitment rather than an empirical conclusion. A second test should be applied to the intervener: What conduct, short of agreement, would show that a person can challenge this allegation in good faith? If the answer is again “none,” the community has closed not only the claim but membership itself.

A closed accusation system is a locked room with a judge inside and no door for new evidence. Facts can be shouted through the window, but every shout is interpreted as another reason to keep the room locked.

5.5 Motive substitution and mind reading

When an argument is difficult to answer, the speaker's motive is substituted for the argument. The critic must secretly admire the villain, share the villain's vice, seek attention, protect status, or act from prejudice. Even if one of these motives were present, it would not show that the argument is unsound.

Motives can matter when assessing testimony, conflicts of interest, or strategic behavior. But motive evidence must itself be evidence. Inferring a hidden vice from disagreement is circular: the critic is corrupt because only a corrupt person would disagree, and disagreement proves the corruption.

Motive substitution is especially corrosive because it creates indefinite allegations. A premise can be examined; a secret character defect can always be redescribed. The conversation moves from public reasons to amateur prosecution of an inner life. Attribution then becomes asymmetric: the crowd's overstatement is explained by righteous concern, while the critic's precision is explained by hidden depravity.

It is easier to invent an X-ray of another person's soul than to answer a difficult sentence. The supposed X-ray cannot be checked, so it can display whatever motive the group needs.

5.6 Crowd authority and compelled conformity

An audience may be told that every informed, decent, or professionally qualified person accepts the charge. The relevant sources remain unnamed, and the consensus is inferred from the local room. Disagreement is therefore treated as proof of ignorance.

Consensus can be powerful evidence when the field, question, methods, independence, and degree of agreement are specified. A crowd in a partisan forum is not that kind of consensus. Nor is an appeal to unnamed experts a substitute for their analysis. The demand to conform to the room is social discipline, not epistemic support.

The phrase "everyone knows" is particularly useful because it obscures the relevant denominator. It can mean everyone in one feed, profession, activist network, congregation, party, or friendship circle. It can also mean that no one present is willing to incur the cost of saying otherwise. Apparent unanimity should therefore prompt questions about selection, independence, private belief, and sanctions for dissent before it is treated as epistemic authority.

The crowd version is often blunt: *"Anyone informed already knows this. Educate yourself before you embarrass yourself again."* The insult hides the missing work. No source is named, no proposition is narrowed, and no path is offered by which an informed person might reasonably disagree.

If everyone at a meeting agrees, first ask who felt welcome at the meeting and what happened to the last person who disagreed. A silent chair can be evidence of absence, but it can also be evidence that someone left.

5.7 The black-sheep penalty

Internal critics are often punished more intensely than external opponents. The opponent is already categorized and can be ignored. The internal critic competes over the meaning of the group's own values. By saying that opposition to the villain does not require the present overstatement, the critic exposes a distinction the loyalty test was designed to erase.

This helps explain the peculiar fury directed at people who agree with nearly the entire indictment but contest one inference. Their proximity is not experienced as reassurance. It makes the residual disagreement more threatening, because the dissenter cannot be dismissed as simply uninformed about the group's moral concern. The dissenter demonstrates that the concern can coexist with a better standard.

The black-sheep penalty converts agreement into an aggravating circumstance. “You know how bad the villain is, so you have no excuse.” The critic’s familiarity with the evidence, prior loyalty, or record of condemning established wrongdoing is used not to assess the objection but to heighten the charge of betrayal. In this way the group protects the inference that only vice could explain dissent.

This resembles a family becoming angrier at the relative who names a household problem than at the outsider who already dislikes the family. The insider’s words carry more danger because they cannot be dismissed as ordinary hostility.

5.8 Personal degradation as argument

Insults can express anger without purporting to reason. The deeper problem arises when degradation is presented as a response to evidence. The target or critic is called stupid, diseased, sexually deviant, intoxicated, delusional, or morally filthy. The abundance of contempt is expected to create the impression that a rebuttal has occurred.

No conclusion follows. The insult may strengthen group solidarity, reward the speaker, and warn observers against dissent. Its function is not to demonstrate falsity but to make continued argument humiliating. It also supplies a low-cost way for marginal members to prove conformity. A person need not understand the contested evidence to repeat the group’s degradation at the approved intensity.

The method throws rhetorical dung at the speaker to see what sticks: “*You are a fraud, a creep, and too dull to understand what decent people can see,*” or “*Your polished language cannot hide the rot underneath.*” The labels are deliberately broad. If one misses, another may still leave a stain in the audience’s memory.

Throwing mud can make a person look dirty, but it does not tell us what the person did. The visual effect is the point. Personal abuse works the same way when its dirtiness is mistaken for proof.

5.9 Punishment before specification

The social-closure sequence often ends with demands for removal, firing, exclusion, investigation, loss of platform, or other punishment before the alleged violation is precisely identified. The proposed sanction supplies emotional gravity to a case that lacks argumentative structure.

Some speech and conduct warrant institutional consequences. But a defensible demand should identify the act, the rule, the evidence, the responsible authority, and the proportionate sanction. “This person should be punished because this person is bad” is the institutional form of the villain exception.

Otherwise the process resembles announcing a sentence before reading the charge. The punishment supplies drama while the missing specification goes unnoticed.

5.10 A field guide to the vitriolic scripts

The following lines are invented composites. They borrow the argumentative shape and emotional heat of common exchanges, not the wording, people, or allegations in any particular dispute. Every example can be reversed across political factions.

1. **The loyalty oath.** “*If you will not repeat the charge, you have shown which side you are on.*” The line does not ask whether the charge is true. It asks the listener to wear it like a team jersey. Refusal is then treated as defection, even if the listener condemns the target and merely doubts this one claim. The tactic works by making a sentence serve two jobs at once: a factual assertion and a badge of membership. A useful reply is, “Which part of the charge is supported, and why must group loyalty depend on repeating the rest?”

2. **Decency in place of proof.** *“No decent person would spend this much time asking for proof.”* Time spent checking the claim is redescribed as a moral defect. Yet serious or complicated accusations often require more work, not less. A bridge inspector does not become hostile to public safety by examining a suspicious crack carefully. The script moves the burden away from the person making the claim and onto the character of the person checking it.

3. **Procedure as secret allegiance.** *“Your concern for procedure tells us whom you secretly value.”* A public rule is converted into a private motive. Because fair procedure may help an accused person, the speaker assumes that anyone defending the procedure must value the accused more than the harmed. But brakes benefit reckless drivers as well as careful ones; that does not make traffic safety a secret loyalty program. The proper question is whether the procedure is fair, workable, and applied consistently.

4. **The beneficiary fallacy.** *“You are not correcting a claim; you are giving cover to garbage people.”* The correction is judged by who may benefit rather than by whether it is accurate. This erases the difference between intending to protect a wrongdoer and refusing to add a falsehood to a true indictment. Correcting a scoreboard can help a team that cheated earlier, but the earlier cheating does not make the wrong score right. A benefit may justify examining motive; it cannot determine truth.

5. **The counterfeit consensus.** *“Look around. Everyone who understands the issue settled this already.”* A visible local majority is presented as if it were independent expert agreement. The line hides questions about who entered the room, who left, who stayed silent, whether the speakers copied one another, and what evidence they actually saw. Ten mirrors can display ten images while reflecting one object. The size of the chorus matters only after its independence and access to evidence are known.

6. **Humiliation in place of a citation.** *“If you were half as intelligent as you pretend, you would not need a citation.”* The insult supplies heat at exactly the point where a source should supply light. It also warns observers that asking the same question may cost them status. A librarian who mocks a reader still has not located the book. The simple repair is to ignore the intelligence test and ask again for the source, the page, and the proposition it is supposed to support.

7. **The self-sealing association.** *“The harder you deny the association, the more obvious it becomes.”* Denial is built into the accusation as confirming evidence, so the target has no possible reply that counts against it. Silence can be called evasive, calm denial rehearsed, and angry denial revealing. This is a locked room with a judge inside and no door for new evidence. A claim that predicts every response distinguishes nothing and should lose confidence, not gain it.

8. **The compassion trap.** *“Either accept the accusation or admit that the harmed people mean nothing to you.”* The script forces a choice between instant belief and moral indifference. It erases positions such as protecting complainants, preserving evidence, investigating quickly, and withholding a final verdict until the facts are clearer. Evacuating after a fire alarm can be wise without treating the alarm as proof that the building burned. Care can guide precaution without dictating every factual conclusion.

9. **Condemnation without a charge.** *“I do not care which offense it was; any version tells us what the target is.”* The speaker openly abandons the task of identifying what happened while preserving the moral verdict. Different acts have different victims, evidence, defenses, and remedies. A blank charge sheet cannot be repaired by writing “bad person” across the top. Once the offense becomes interchangeable, punishment is aimed at an identity rather than a demonstrated act.

10. **Prestige laundering.** *“A respected public figure said it openly, so only a fool still doubts it.”* Publicity and status are treated as if they manufacture evidence. Authority can matter when it brings relevant expertise, access, method, or a record of reliability. Without those links, a famous speaker is only a louder speaker. A celebrated chef may know food, but the chef’s confidence does not settle who caused a traffic collision. The bridge from prestige to knowledge must be shown.

11. **Complexity as confession.** *“Your long explanation is just a polished defense of filth.”* The response is rejected because it is long, qualified, or carefully written. Yet a claim assembled from many accusations may require a long answer simply to separate them. Calling a map evasive because it contains many roads does not tell us which road is wrong. The relevant test is whether the explanation answers the claims, uses sound evidence, and leaves room for correction.

12. **Technical language as moral disguise.** *“People like you hide behind technical language because ordinary morality exposes you.”* Technical language can certainly be used to confuse or intimidate, but it can also protect distinctions that ordinary speech blurs. Legal, medical, and statistical terms often exist because nearby ideas have different consequences. A mechanic naming the failed part is not hiding from the broken car. The fair response asks for a plain-language explanation and then checks whether the distinction is real.

13. **The diagnostic trap.** *“Keep talking. Every sentence shows the room what you really are.”* The speaker declares that the discussion is no longer about reasons. Every answer is now raw material for a diagnosis of character. Continued explanation proves obsession; withdrawal proves guilt; frustration proves exposure. It is the conversational version of a coin that means guilty on both heads and tails. Once no response could count as good faith, the exchange has become a performance of power rather than an inquiry.

14. **The joke ratchet.** *“It was only a joke. Why are you so defensive unless it struck a nerve?”* The statement is serious enough to spread suspicion but unserious when evidence or responsibility is requested. The objection is then recycled as proof that the joke found a hidden truth. This gives the speaker two exits while giving the target none. A rubber sword cannot become steel when it wounds and rubber again when someone asks who swung it.

15. **Punishment first.** *“Someone who argues this way should not be allowed to teach, publish, or hold office.”* A sanction is proposed before the false statement, broken rule, responsible institution, and proportionate remedy are identified. The severity of the punishment then makes the accusation feel more serious, as if a requested sentence supplied evidence for the charge. This is a courtroom that announces the penalty before opening the file. The demand should be paused until the act and rule can be stated plainly.

16. **Expulsion as victory.** *“Stop talking and leave. Decent people are finished listening.”* Removal ends the practical discomfort of answering and then masquerades as a successful rebuttal. The remaining group may mistake the new silence for agreement, although the disagreement was merely pushed outside. Turning off the scoreboard does not win the game. Expulsion may sometimes be needed for threats or disruption, but it should not be confused with answering an argument.

The surface insults vary, but the lines perform the same small set of jobs. They shift attention from the claim to the speaker, treat revulsion as proof, make disagreement diagnostic, and threaten the dissenter’s status. Their ugliness is useful to the tactic. A calm paraphrase can hide how hard such language makes it to think, answer, or remain in the room.

These examples are not evidence that one faction behaves worse than another. They illustrate portable scripts. Change the villain, the vice, and the approved crowd, and the same sentence can travel across the aisle.

Diagnostic map: what the scripts replace

What inquiry needs	What the script substitutes	Immediate social payoff	Long-term cost
A precise claim	A judgment about identity or allegiance	The group quickly knows who is “with us”	The original proposition becomes hard to locate or correct
Claim-specific evidence	Disgust, prestige, repetition, or insult	Condemnation becomes fast and emotionally satisfying	Weak and strong allegations become difficult to distinguish
A possible route to revision	A closed diagnosis in which every response confirms guilt	The speaker never has to concede uncertainty	Error becomes self-protecting and correction becomes suspicious
A proportionate remedy	Expulsion or punishment before specification	The accusation gains drama and coercive force	Institutions act on labels rather than established conduct
Independent agreement	A visible crowd repeating dependent claims	Public confidence looks larger than private confidence	Counterfeit consensus becomes mistaken for corroboration

6. Sociological Inertia and the Irrationality Ratchet

An individual error can disappear when one person changes a mind. A social error can keep moving even after many people privately step away from it. Public statements, damaged friendships, institutional decisions, archived accusations, and expectations about one another remain. We therefore need to explain not only how people reason badly, but how bad reasoning learns to walk on its own.

The key is a feedback loop. One person’s shortcut makes the same shortcut easier for the next person. The second performance seems to validate the first, not by adding evidence, but by making the behavior look normal. Public praise marks it as desirable. Dissent becomes harder to see, so the consensus looks stronger, which makes dissent still more costly. Call this the *irrationality ratchet*: group confidence and rhetorical heat can click upward one notch at a time, while moving back down takes much more work.

A stadium wave shows the basic shape. Each person stands because nearby people stood. From a distance the wave looks like one organized decision, but no central signal may exist. In public judgment, however, observers often mistake the size of the wave for the strength of the evidence.

The irrationality ratchet at a glance

1. Shortcut	2. Repetition	3. Apparent confirmation	4. Dissent penalty
A claim outruns its support because the target already seems bad	Others copy the conclusion, label, or level of certainty	Dependent repetition looks like independent evidence and private doubt stays hidden	Correction threatens identity, reputation, and prior sanctions, so dissenters are assigned the villain’s vices

The fourth stage feeds the first. Once correction is socially costly, the next unsupported claim meets fewer visible obstacles and begins the circuit with more borrowed confidence.

6.1 Descriptive validation and injunctive validation

Descriptive validation answers, “Do people like us do this?” Injunctive validation answers, “Will people like us approve if I do this?” The two forms of validation are analytically distinct even when one gesture supplies both. A widely shared accusatory slogan displays frequency; the rewards attached to sharing it display approval. A correction that receives silence displays rarity; the abuse directed at its author displays disapproval.

This matters because the observer is not necessarily choosing between rational independence and conscious dishonesty. The observer may treat the crowd as information. If many respected peers state a conclusion confidently, it is reasonable in ordinary life to suppose they possess relevant knowledge. Social learning is indispensable; no person can verify everything independently. The failure begins when the observer cannot distinguish a chain of copied confidence from multiple independent inquiries.

The same crowd also changes incentives. A person who privately assigns a claim only moderate probability may speak categorically because qualified language has been marked as suspect. The categorical statement is then observed by others who infer categorical private confidence. Public certainty exceeds private certainty, but every participant experiences the excess as evidence about everyone else.

Thus the social validation of irrationality is not merely comfort after an error. It can alter the perceived evidence before the next judgment. The fact that others use the same shortcut reduces shame, supplies an apparent method, and produces a misleading signal of corroboration. A weak inference becomes psychologically easier and sociologically safer at the same time.

An applause meter captures both sides. Loud applause tells a speaker what the room likes, but later listeners may also hear the applause as proof that the speaker was right. Approval has been mistaken for verification.

6.2 Matching the type and degree of irrationality

Groups do not validate only conclusions. They validate styles and magnitudes of departure from rational discipline. Members learn whether the approved response is to omit uncertainty, reverse the burden, infer motives, repeat a rumor, use a diagnostic label, demand punishment, or degrade the critic. They also learn how far to go.

This creates an expressive escalation problem. Once ordinary condemnation becomes ubiquitous, it ceases to distinguish especially committed members. Stronger labels, broader allegations, greater certainty, and harsher sanctions can restore the signal. The baseline moves. Yesterday's maximal denunciation becomes today's minimum demonstration of seriousness.

Matching intensity has a defensive purpose. An understated response can be interpreted as concealed doubt, and concealed doubt as concealed vice. The safest strategy is therefore to reproduce the local package: the same label cluster, the same dismissive account of counterevidence, the same personal conclusions about dissenters, and the same degree of moral confidence. Individual performances converge even when the underlying reasons do not.

The process can become an arms race. If saying "wrong" no longer proves commitment, members say "evil." When "evil" becomes ordinary, someone reaches for "monstrous." The escalation may reveal more about the competition for standing than about any new fact.

Online feedback can accelerate this learning. Brady, McLoughlin, Doan, and Crockett (2021) found that positive social feedback increased the likelihood of future outrage expression and that users learned outrage norms from their networks. The point is not that online outrage is generally irrational. It is that reinforcement can train the expression independently of any improvement in claim-specific evidence. A community can thereby learn the preferred intensity of condemnation faster than it learns whether the allegation is true.

6.3 Thresholds, cascades, and path dependence

Granovetter's (1978) threshold model shows how collective outcomes can depend on the proportion of others already participating. People have different thresholds for joining a behavior: one acts immediately,

another after a few have acted, and another only after a majority. Small differences in the distribution or sequence can therefore produce large differences in the final crowd even when groups have similar average dispositions.

Applied to villain discourse, early accusers and high-certainty speakers can trigger people with progressively higher thresholds. The visible crowd then appears to reflect uniformly strong private conviction, although it may contain zealots, cautious conformists, strategic actors, and people merely unwilling to become the next target. Aggregate extremity is not a transparent window into individual belief.

Bikhchandani, Hirshleifer, and Welch (1992) modeled informational cascades in which later actors rationally follow earlier actions while discounting their own private information. Their formal setting differs from moral and political discourse, but the insight transfers carefully: when actions are visible and the information behind them is not, people can infer too much from the sequence. A person may repeat an accusation because others repeated it, while later observers interpret the repetition as independent confirmation.

The resulting path dependence explains sociological inertia. If the first visible judgments had been cautious, later thresholds might have produced a cautious norm. Once a categorical cascade has formed, however, contrary private information must overcome both the original claim and the social evidence manufactured by the cascade. Correction is no longer a simple update. It is a coordination problem.

Dry grass offers another analogy. One spark may die in one field and race across another because the surrounding conditions differ. Early confident speakers are sparks; low social thresholds are dry grass. The size of the later fire does not tell us how strong the first spark was.

6.4 Pluralistic ignorance and the counterfeit consensus

A group can publicly sustain a position that many members privately regard with doubt. Prentice and Miller (1993) documented pluralistic ignorance in a different domain: students misperceived peers' comfort with campus drinking, and some shifted over time toward what they mistakenly believed the norm to be. The relevance here is structural. People infer private attitudes from public behavior even when that behavior was itself shaped by perceived attitudes.

Villain-centered communities are especially vulnerable because restraint has an identity cost. A privately doubtful member sees categorical condemnation from others, not the doubts those others withheld. The member concludes that uncertainty is deviant and either stays silent or performs certainty. Each defensive performance makes the original misperception more plausible.

This produces a counterfeit consensus. It is real as public behavior but unreliable as a measure of private judgment or independent evidence. Institutions may nevertheless cite it: editors invoke public expectations, employers invoke organizational values, politicians invoke constituent demands, and platforms invoke community norms. The accumulated references make a fragile pattern appear settled.

Pluralistic ignorance also explains why one dissenting voice can matter disproportionately. The dissenter does more than offer an argument. The dissenter reveals that doubt can be spoken and survived. That revelation may lower others' expressive thresholds. Vice transfer and the black-sheep penalty prevent this cascade from running in reverse.

The pattern resembles the story of the emperor's clothes. Many observers may notice the problem, yet each sees only everyone else's public silence. One plain statement can matter because it reveals that private sight was never as rare as it appeared.

6.5 Repetition, source laundering, and borrowed certainty

The irrationality ratchet depends on losing track of provenance. A single allegation passes through news summaries, screenshots, commentary, memes, institutional statements, and conversations. At each stage the reporting verb weakens. “One person alleged” becomes “reports indicate,” then “the controversy involves,” then “everyone knows.”

This is source laundering: dependency is concealed by changes in format and speaker. The same claim acquires several reputational wrappers without acquiring another evidential foundation. A later speaker borrows certainty from the number and status of prior speakers; prior speakers then cite the later repetition as proof that the claim has become widely established.

Think of a rumor sent through a hall of mirrors. Each reflection looks like another object, but all of them depend on the same original. Changing the frame from a post to a headline to an official statement does not by itself create a new witness.

The correction travels at a disadvantage. It must restore the original proposition, trace the chain, distinguish primary from secondary sources, and explain why apparent plurality is not independence. The allegation needs only to fit the villain narrative. The asymmetry gives error a lower transmission cost than repair.

Even accurate allegations can be damaged by this process. Source laundering makes it difficult to distinguish the well-supported core from embellishments added during transmission. When one embellishment collapses, defenders may protect it to preserve the global narrative, while opponents use it to discredit the entire case. Sociological validation has converted an evidentially mixed record into an all-or-nothing identity contest.

6.6 Why the ratchet resists reversal

Several investments accumulate as the cascade grows. Individuals have made categorical statements. Organizations have issued sanctions. Journalists and scholars have attached reputations to interpretations. Friends have punished friends. To revise is no longer only to change a belief; it is to admit that prior confidence, aggression, or punishment exceeded the evidence.

At that point a correction feels less like replacing one loose brick and more like admitting that part of the house was built on it. People may defend the brick because they fear what else must be repaired.

The second villain absorbs this pressure. Instead of revisiting the claim, participants explain correction as infiltration, moral weakness, corruption, or betrayal. The intervener’s alleged character becomes a reason not to inspect the intervention. More intervention produces more vice attribution, which appears to confirm that vigilance was necessary.

Group polarization can raise the reversal cost further. Discussion among like-minded participants may generate new arguments, competitive demonstrations of commitment, and a more extreme comparison standard. Once the group has defined itself partly through maximal opposition to the villain, reducing confidence in any adjacent allegation can feel like reducing the group’s moral identity.

The ratchet metaphor should not imply irreversibility. Cascades can be fragile, norms can change, and well-placed dissent can reveal hidden agreement. But reversal requires more than one correction. It requires changing the social meaning of correction: from aid to the villain into maintenance of the group’s own epistemic integrity. Without that change, every private doubt must independently bear a public cost, and sociological inertia will outlast the evidence that initiated it.

7. Rhetorical Dung and Connotative Scattershot

7.1 The label barrage

One of the most recognizable products of villain thinking is the tendency to throw rhetorical dung: a rapid series of terms with extreme negative connotations, offered not as carefully supported classifications but as contaminating matter. The tactic can be called *connotative scattershot*. The speaker fires many morally explosive labels in the hope that at least one will stick, or that the combined residue will make proof of any one label unnecessary.

A political opponent may be called, in quick succession, authoritarian, fascist, communist, traitorous, terrorist, racist, genocidal, cultic, criminal, predatory, corrupt, or insane. The list is politically reversible. Different coalitions prefer different projectiles, but the logic is the same.

These terms are not synonyms. They have different historical, legal, psychological, and political criteria. A person can be authoritarian without being fascist, prejudiced without being genocidal, criminal without being a traitor, sexually abusive without meeting the clinical meaning of a paraphilic disorder, or radically ideological without belonging to a cult. Sometimes several labels are independently warranted. Connotative scattershot avoids the work of showing which ones.

The tactic is like throwing a bucket of mixed paint at a wall and then claiming that every separate color has been carefully applied. The wall is certainly stained. That does not tell us which label fits.

7.2 Emotional convergence is not evidential convergence

The labels feel mutually supportive because their connotations converge. Each says, at a high level of abstraction, “bad and dangerous.” But evidential support does not converge merely because emotional valence does.

Consider five allegations supported by five different weak analogies. Combining them may create an intense global impression while leaving each classification unsupported. The audience mistakes convergence of disgust for convergence of proof.

Five smoke machines can make a room look dramatic without producing a fire. In the same way, several words that all evoke horror can create one strong mood without creating one strong case.

The same process can reverse the ordinary relation between labels and evidence. In rational argument, evidence supports a description, and the description helps determine an appropriate moral response. In connotative scattershot, the desired moral response selects the harshest descriptions, and those descriptions are then treated as evidence for one another.

7.3 Definition shopping

When one label is challenged, the speaker searches for an unusually broad definition under which it might apply. If the target has intimidated anyone, perhaps “terrorist” can be stretched to fit. If the target favors strong executive action, perhaps “fascist” can be defined as any authoritarian tendency. If the target repeated a falsehood, perhaps “fraud” can be detached from its legal elements. If the target supported a harmful policy, perhaps “murderer” can be extended through a long causal chain.

Broad stipulative definitions are not always illegitimate. Concepts can have ordinary, technical, legal, and rhetorical senses. The abuse occurs when a broad sense is used to secure application while the term’s narrower

and more catastrophic connotations are used to secure condemnation. The speaker wins extension with one meaning and emotional force with another.

It is a bait-and-switch with definitions. The wide definition gets the target through the door; the narrow term's terrible history supplies the punishment once inside.

7.4 Contamination without commitment

Scattershot gives the speaker deniability. If one term is refuted, the others remain in circulation. If all are challenged, the speaker can say the list was expressive rather than literal. Yet audiences retain the association. The target has been placed in a semantic cloud of atrocity.

This creates contamination without commitment. The speaker receives the reputational effect of accusation while avoiding responsibility for a proposition precise enough to test. The tactic is therefore not only abusive language. It is a method of maintaining maximum condemnatory force with minimum falsifiability.

The stain stays even when the thrower says, "I did not mean any one word literally." The retreat removes the speaker's burden more easily than it removes the audience's association.

7.5 Platform incentives

Online environments reward moral-emotional language. Brady and colleagues (2017) found that moral-emotional terms increased the diffusion of political messages within ideological networks. Brady, McLoughlin, Doan, and Crockett (2021) further found that social feedback and norm learning can amplify expressions of moral outrage. The ecology of attention therefore favors the vivid label over the qualified description.

The individual may learn that "the evidence suggests a serious but unresolved possibility" receives little response, while a maximal accusation earns likes, shares, visibility, and group membership. With repetition, overstatement becomes not merely a tactic but the learned sound of sincerity. Moderation feels false because the community has trained itself to hear outrage as proof of concern.

7.6 Label intensity as a membership metric

Once a label cluster becomes conventional, its content and its social use can separate. Members may repeat words without agreeing on criteria because the shared vocabulary certifies that they know who the enemy is. The terms become passwords. Their semantic looseness is an advantage: different participants can project different wrongs into one common denunciation.

This helps explain why correcting one label can provoke more hostility than rejecting a factual detail. The factual correction threatens one proposition. The label correction threatens a membership ritual. Asking whether the target meets the criteria for a loaded term is heard as refusing the group's linguistic uniform.

The labels work like team scarves. Wearing the scarf shows where one belongs even if the wearer cannot explain the team's strategy. Refusing the scarf is then treated as switching sides.

The degree of rhetorical excess can then become a metric of sincerity. A speaker who uses one earned term appears less committed than a speaker who uses five maximal terms. A member who concedes uncertainty appears morally behind a member who reports certainty. Scattershot is socially selected because it displays that no residual sympathy or hesitation remains.

Vice transfer reinforces the metric. The critic who refuses the label barrage is not merely described as cautious; the unused labels are redirected toward the critic. Refusal to throw the dung is treated as evidence that one belongs in its path. The tactic thereby recruits future throwers through both reward and warning.

7.7 The connotative audit

A useful corrective is to audit each loaded term:

1. What descriptive criteria does the term assert?
2. Which items of evidence support each criterion?
3. Is the term being used in an ordinary, technical, clinical, historical, or legal sense?
4. Would the same criteria be applied to an ally?
5. What less inflammatory description captures exactly what the evidence establishes?
6. What information would show that the term does not apply?

The purpose is not to ban severe words. It is to make severity answerable to description. Some people are fascists, some acts are terroristic, some policies are racist, some campaigns are genocidal, and some officials are corrupt. Precision does not weaken justified condemnation. It distinguishes it from dung-flinging.

A sharp knife is more useful than a bucket of broken glass. Precise moral language can cut along real differences. A random spray of harsh terms merely wounds the discussion.

8. What the Villain Exception Does to Justice

8.1 Justice becomes target-indexed

Justice requires more than reaching a correct result by accident. It requires procedures that can distinguish correct from incorrect results across cases. When evidential burdens change with the identity of the accused, judgment becomes target-indexed: the rule is not “apply this standard to this kind of claim” but “apply the standard needed to condemn this person.”

The temptation is understandable. A villain may have escaped accountability before. Institutions may be slow, compromised, or unequal. Demanding ordinary process can seem like granting another advantage to the powerful. Yet a procedure that becomes optional when the audience is certain of guilt is no procedure at all. Certainty is precisely the state in which bias becomes least visible to the believer.

A ruler that changes length depending on what it measures is not a ruler. A court, newsroom, or workplace process that changes its burden according to whom it dislikes has the same defect.

8.2 Standards of proof are not statements of sympathy

Different institutions properly use different thresholds. Informal belief, journalistic publication, workplace precaution, civil liability, and criminal punishment need not require the same degree or kind of evidence. Confusing them produces two opposite errors: demanding criminal-proof standards before taking any protective action, or treating the low threshold for investigation as though it established guilt.

A disciplined vocabulary helps:

- *reported*: a source has made the claim;
- *alleged*: the claim has been asserted but not established;
- *corroborated*: independent evidence supports material parts of the claim;

- *substantiated*: an inquiry found adequate support under a stated standard;
- *adjudicated*: an authorized tribunal reached a finding;
- *liable*: civil responsibility was found under the governing civil standard;
- *convicted*: criminal guilt was found under the governing criminal standard.

These categories can overlap, and institutions vary, but they should not be silently collapsed. Legal and clinical terms are not emotional superlatives. Misclassification is not made noble by the ugliness of the underlying conduct.

The categories are steps on a staircase, not different names for the top floor. A report may justify an investigation without justifying a conviction. Confusing the steps lets a person ride the emotional elevator from accusation to guilt.

8.3 Punishment becomes fungible

If the target deserves punishment for offense A, audiences may become indifferent to whether punishment is formally imposed for unsupported offense B. The total suffering feels deserved even if the stated basis is false.

This is a profound failure of justice. Punishment is not a general debit against a bad person. It is an institutional response to specified conduct under specified authority. Fungible punishment licenses fabricated charges against people who are genuinely guilty of something else, a practice that destroys the distinction between accountability and vengeance.

Punishment is not store credit that can be charged against any item on a bad person's account. A person who deserves a penalty for A does not thereby deserve a false record saying B.

8.4 Victims become instruments

The villain exception is often defended in the name of victims, yet it can instrumentalize them. A complainant's suffering becomes a shield against questions the claim still requires. The audience may elevate a person as symbol and then discard the person's actual account when it no longer serves the preferred narrative.

Serious treatment of victims requires preservation of testimony, competent investigation, protection from retaliation, trauma-informed procedure, and honest communication about what is known. It does not require declaring every allegation infallible. Credulity and hostility are not the only options.

A victim should not be turned into a team mascot whose story is simplified for the crowd. Respect means listening closely enough to preserve the person's actual account, including its limits and uncertainties.

8.5 Outrage can distort adjudicative confidence

Moral emotion appropriately informs judgment, but it can also alter the experience of evidential sufficiency. Salerno and Peter-Hagene (2013) found that the co-occurrence of anger and disgust predicted moral outrage and, in their mock-jury study, greater confidence in a guilty verdict. The relevance is not that jurors or citizens should feel nothing. It is that the feeling of certainty can be partly produced by moral emotion rather than by additional proof.

The rational response to strong emotion is not automatic distrust of one's judgment. It is procedural supplementation: slow the inference, separate charges, specify standards, invite adversarial testing, and distinguish the strength of one's moral reaction from the strength of the evidence.

Emotion is fuel. It can move people to investigate, protect, and repair. It should not also be the steering wheel, the map, and the proof that the destination has been reached.

8.6 The procedural defender becomes an accused party

Justice is impaired when the person defending a procedure must first prove moral distance from the accused. Judges, lawyers, editors, investigators, witnesses, scholars, and ordinary citizens occupy roles whose public value depends on conclusions not being assigned by coalition. If every favorable ruling, correction, or evidential qualification is treated as evidence that the role-holder shares the defendant's vice, impartial work becomes reputationally hazardous.

This pressure can alter decisions without any explicit command. An investigator may broaden a conclusion to avoid appearing sympathetic. An editor may retain an unsupported insinuation to avoid seeming protective. An institution may impose a sanction whose public symbolism exceeds its findings. The target-indexed standard has migrated from public argument into the incentives of adjudication.

Role protection does not confer immunity. A lawyer can suborn perjury, an editor can suppress relevant evidence, and an investigator can be corrupt. But those are independent allegations requiring independent support. It is not legitimate to infer them from the fact that a role was performed in a way that incidentally benefited a villainized party.

When procedural defenders become accused parties, the quality of justice falls for everyone. Competent people avoid the role, remaining participants learn to anticipate the crowd, and the appearance of consensus grows. The justice system can still retain its formal steps while losing the social conditions under which those steps produce independent judgment.

The officials begin to act like weather vanes. Nothing in the written procedure orders them to follow the wind, yet they learn which direction is safest. The form remains while independent judgment drains away.

8.7 Justice must preserve proposition ownership

Every finding should retain an identifiable owner, scope, and standard. A witness owns a testimony, an investigator owns an investigative conclusion, a tribunal owns a judgment, and the public may form further opinions. Trouble begins when one layer borrows the authority of another. A social-media allegation is described as though a court established it; an institutional precaution is described as a factual finding; a crowd's moral verdict is described as professional consensus.

Proposition ownership also protects correction. The body that made a claim should be able to revise that claim without every adjacent judgment collapsing. If allegations, findings, and moral evaluations have been fused, any correction threatens the whole structure and will be resisted accordingly.

Justice therefore requires a public record that identifies who concluded what, from which evidence, under which standard, with what remaining uncertainty. This discipline may look bureaucratic beside the immediacy of outrage. It is precisely what prevents sociological momentum from laundering accusation into adjudication.

Think of labels on folders. "Complaint," "investigator's finding," and "court judgment" belong on different folders even when they concern the same event. Pouring every page into one folder marked "guilty" makes later checking almost impossible.

9. What the Villain Exception Does to Public Discourse

9.1 It makes precision look immoral

Public reasoning depends on people being allowed to make distinctions without having every distinction interpreted as exoneration. The villain exception removes this permission. To distinguish accusation from adjudication, association from participation, one offense from another, or possibility from probability is to risk being placed on the villain's side.

The predictable result is not better agreement. It is less truthful speech. Participants either repeat the maximal claim, leave the discussion, or retreat into private doubt. The visible conversation becomes more certain as the underlying audience becomes less candid.

Precision is the map legend of public discussion. If every shade is forced into "good" or "evil," the map becomes vivid but useless for finding where anything actually is.

9.2 It replaces proposition exchange with status display

Argument should expose propositions to reasons. Villain discourse instead asks speakers to display the proper moral posture. The most valued contribution is not the best-supported claim but the strongest sign of opposition. Condemnation becomes positional: every participant must stand farther from the villain than the last.

This dynamic produces escalation. If one person says the target is dishonest, another says criminal; if one says authoritarian, another says fascist; if one says dangerous, another says genocidal. The labels establish rank within the condemning coalition. Evidential restraint becomes a status loss.

The discussion becomes an auction in which people bid with stronger words. The winner is the person who sounds most opposed, not the person who offers the best-supported description.

9.3 It destroys correction

Correction is difficult when a claim is attached to identity. Withdrawing an accusation can feel like helping the villain, betraying victims, or validating the opposing tribe. The correction is therefore narrowed, buried, or made conditional. A specific falsehood is conceded while the speaker insists that the target's general badness preserves the original point.

But the general point was not the only thing communicated. A specific accusation entered public memory. Unless correction is comparably visible and explicit, the false detail continues to shape belief through familiarity. The public archive accumulates emotionally useful claims after their evidential basis has evaporated.

A correction whispered after an accusation was shouted is not equal repair. It is like removing a false sign after thousands of travelers have already followed it and then leaving no notice that the route was wrong.

9.4 It creates false consensus and false extremity

Online outrage can distort observers' beliefs about groups. Brady and colleagues (2023) found that users overperceived the moral outrage expressed by others online, and experimental manipulations of this overperception inflated beliefs about hostile communication norms, affective polarization, and ideological extremity. Moore-Berg, Ankori-Karlinsky, Hameiri, and Bruneau (2020) likewise found that exaggerated beliefs about how negatively opposing partisans viewed one's own group predicted intergroup hostility.

The villain exception feeds these meta-perceptions. Each side sees the other's harshest accusations, infers a group norm of hatred, and licenses retaliatory harshness. Public discourse becomes a hall of mirrors in which exaggerated condemnation is both evidence of enemy villainy and justification for one's own escalation.

9.5 It rewards performative certainty

Because visible confidence is often interpreted as knowledge and moral seriousness, uncertainty has poor rhetorical economics. A calibrated statement travels slowly. A categorical accusation is portable, memorable, and identity-confirming.

This does not mean that every confident speaker is performing insincerely. The deeper problem is selection. Networks reward and recirculate the most decisive formulations. Over time, audiences encounter a sample skewed toward extremity and mistake it for the distribution of informed judgment.

It is like judging a town by the people holding megaphones in its square. The loudest voices are real residents, but they are not a random sample of the town.

9.6 It eliminates the internal error-correction function

Healthy discourse needs people who can say, from inside a cause, "our conclusion may be right, but this reason is wrong." The villain exception makes that sentence difficult to utter. Because critics are sorted by whose position benefits, an internal correction is reclassified as an external attack.

The loss is greater than one silenced opinion. Internal critics know the group's sources, language, incentives, and blind spots. They can often identify an error before opponents do and can present the correction in terms the community understands. Treating them as derivative villains removes the least costly form of quality control.

Errors then survive until exposed by an adversary. By that point the claim is attached to group identity, and the correction arrives from a distrusted source. Members defend what they might once have revised. The punishment of the rational intervener thus manufactures the later polarization used to prove that internal discipline was impossible.

An internal critic is a smoke detector inside the building. Disabling it because its noise is embarrassing does not make the building safer. It only ensures that the first warning will come from people outside who may already want the building to fail.

9.7 It makes interlocutors disappear behind inherited vices

Public discourse presupposes that a speaker can be more than the side temporarily benefited by a sentence. Vice transfer denies that presupposition. The critic's evidence, qualifications, and record are replaced by the moral inventory of the villain. The person is no longer heard as an interlocutor but as a carrier.

This collapse changes the grammar of disagreement. Replies no longer take the form "that premise is false" or "that inference is invalid." They take the form "people like you always..." The referent expands from a claim to a type, and the type is defined by the very disagreement under review.

Once inherited vice becomes legitimate argument, reason-giving is optional. A community can defeat any objection by redescribing its author. The public sphere retains the outward appearance of debate - many words, many replies, intense engagement - while proposition exchange has ended.

The intervener is made to wear the villain's shadow. People stop asking what the intervener said because they believe the shadow has already answered.

9.8 It leaves a durable residue after correction

False or exaggerated accusations do not vanish when corrected. Search results, screenshots, summaries, casual memories, and institutional records preserve the earlier formulation. The emotional association may remain even for people who no longer endorse the literal proposition.

Connotative scattershot increases this residue because a retraction usually addresses one term while the semantic cloud survives. Vice transfer broadens the residue to the critic: the very existence of a dispute may be remembered as evidence that both the villain and the intervener were morally suspect.

Public correction must therefore address reach as well as wording. A quiet amendment beneath a viral accusation does not restore the prior epistemic state. Without comparable visibility and clear proposition ownership, discourse accumulates moral sediment that later speakers mistake for history.

The first accusation is wet ink. Even after the sentence is crossed out, the stain may remain on every copy made before the correction.

10. Potential Negative Effects on Society

10.1 Reciprocal norm collapse

The villain exception is rarely contained within one coalition. Each side observes the other's shortcuts and treats them as permission for its own. An unsupported accusation by one faction becomes precedent for another. Selective censorship justifies retaliatory censorship. Politicized prosecution justifies counter-prosecution. Refusal to correct one falsehood becomes a reason not to correct the next.

This is how a norm can disappear while nearly everyone continues to endorse it verbally. Each group claims to favor evidence, fairness, free inquiry, and due process in general, but suspends each principle in response to the other's villainy. The principle survives only as a weapon against opponents.

Reciprocity can sustain good norms when people reward compliance. It can also accelerate collapse when violations are treated as licenses. A society of permanent exception eventually has no standard case left.

Each side keeps a folder of the other's violations and uses it as a stack of permission slips. The pile grows until almost any new shortcut can be excused by an older one.

10.2 Democratic safeguards become negotiable

Polarized citizens may trade institutional principles for partisan advantage. Graham and Svolik (2020), using a nationally representative candidate-choice experiment and a natural experiment, found that only a limited fraction of voters prioritized democratic principles when those principles conflicted with partisan and policy preferences; stronger polarization made this tradeoff worse.

The villain exception supplies the moral narrative for the trade. If the opponent is not merely wrong but existentially evil, procedural restraint appears suicidal. Emergency powers, censorship, selective law enforcement, electoral manipulation, or institutional retaliation can be represented as necessary defenses of democracy. The same measures become proof of tyranny when used by the other side.

Emergency powers are like an axe kept behind glass. Sometimes the danger justifies breaking the glass. But a society that sees every opponent as a fire will soon live among broken cases and swinging axes.

The social danger does not depend on both factions exercising equal power. The faction with greater institutional capacity can do greater harm. The psychological mechanism matters because it helps citizens bless abuses by their own side and interpret restraints as aid to the enemy.

10.3 Degradation of investigative institutions

Journalism, universities, courts, legislatures, and professional bodies depend on distinctions the villain exception erodes. Investigators must separate leads from findings. Editors must distinguish reportability from truth. Courts must distinguish suspicion from admissible proof. Scholars must distinguish hypothesis from result. Institutions lose credibility when they adopt activist confidence without institutional verification.

The damage is double. Partisans first pressure institutions to ratify maximal claims. When correction later arrives, the opposing faction uses the error to dismiss sound work from the same institution. Every evidential shortcut creates material for generalized distrust.

Institutions should not avoid difficult conclusions for fear of controversy. They should make their warrant visible: source quality, uncertainty, standard of review, conflicts, methods, and correction procedures. Credibility is not produced by agreeing with the morally favored audience. It is produced by procedures that remain legible when conclusions are unwelcome.

A bridge inspector earns trust by showing the measurements, not by telling the crowd what it hopes to hear. Investigative institutions need the same visible workmanship.

10.4 Dilution of genuine wrongdoing

Label inflation raises the rhetorical volume while reducing discriminatory power. If every coercive policy is fascism, every grave injustice genocide, every deception fraud, every harmful association conspiracy, and every partisan a traitor, language can no longer mark important differences in degree and kind.

This is not merely a semantic concern. Genuine atrocities require accurate recognition. When maximal terms are used as ordinary abuse, audiences become numb, opponents dismiss the vocabulary as partisan noise, and perpetrators gain room to portray precise accusations as another episode of exaggeration.

The same dilution affects victims. A well-corroborated account is forced to compete in an environment saturated with rumors that use the same language. The public becomes either indiscriminately credulous or indiscriminately skeptical. Both reactions benefit sophisticated wrongdoers.

A fire alarm that sounds for burnt toast, a hot shower, and a burning roof will eventually be ignored. Inflated moral language can create the same dangerous numbness.

10.5 Epistemic segregation

The loyalty test makes shared inquiry costly. People migrate toward spaces in which their doubts will not be interpreted as moral defects. Communities then sort not only by conclusion but by which facts may be questioned. Rival information systems emerge, each with its own villains, authorities, forbidden distinctions, and archive of uncorrected accusations.

Once segregation is advanced, a true statement can be rejected because it is associated with the wrong side, while a false statement can be accepted because rejecting it would aid the enemy. Information is evaluated as a coalition signal before it is evaluated as information.

The factions begin using different maps. A road shown clearly on one map is denied because the other side printed it, while a missing bridge is accepted because it appears on one's own.

10.6 Governance by enemy image

Policy becomes distorted when the central objective is defeating the villain rather than solving the problem. A proposal is assessed by who supports it; compromise is betrayal; acknowledging a useful idea from the opponent threatens the narrative of total corruption. Leaders gain incentives to preserve the villain image because it organizes supporters and excuses failure.

The opponent may indeed pose a serious threat. But a threat analysis should specify capacities, incentives, constraints, and plausible actions. A villain narrative collapses these into essence. Essential enemies cannot bargain, change, respond to incentives, or contain internal differences. Policy built on that representation loses access to ordinary tools of prediction.

A caricature may help a crowd recognize whom to boo, but it is a poor map for governing. Policy needs details that the caricature has erased.

10.7 Escalation from speech to coercion

Connotative scattershot lowers the psychological barrier to coercion. A person described as mistaken should be answered. A person described as corrupt may be investigated. A person described as an existential predator, traitor, terrorist, or subhuman contaminant appears to require removal.

Not every harsh label produces violence, and strong language can be justified. The risk lies in cumulative moral exclusion. Dehumanization, belief in pure evil, and moral typecasting each reduce concern for what may be done to the target (Gray & Wegner, 2009; Kteily et al., 2015; Webster & Saucier, 2013). When the target is no longer perceived as a possible victim of injustice, punishment can expand without a corresponding demand for proof.

Words can paint a target red before any order is given. Once the person is marked as outside the circle of protection, harsher treatment feels less like a new decision and more like following the label.

10.8 Epistemic nihilism

Repeated exposure to accusation, counter-accusation, correction, and retaliatory fabrication can produce exhaustion. Citizens stop asking which claims are true and conclude that all claims are propaganda. This cynicism may look sophisticated, but it abandons discrimination as completely as credulity does.

The villain exception contributes to this result by making accuracy visibly subordinate to allegiance. When audiences observe that each faction's evidential standards reverse with the target, they learn that public reason is theater. The long-term casualty is not only belief in a particular fact. It is belief that facts can govern political judgment at all.

Citizens come to see fact-checking as stage scenery: moved into place when useful and carried away when the scene changes. The result is not healthy doubt but surrender.

10.9 Institutional path dependence

Sociological inertia becomes materially consequential when institutions act on a cascade. A preliminary accusation may trigger a public statement, staffing change, policy revision, funding decision, disciplinary file, or database entry. Later participants encounter the institutional response and infer that a verified finding must exist behind it.

The inference can run in a circle. The institution acted because public confidence was intense; the public becomes more confident because the institution acted. Each side cites the other as independent validation even

when both ultimately rely on the same uncertain claim. Reversal then threatens not just reputations but budgets, careers, and legal exposure.

Institutional memory can outlast the individuals who formed the judgment. New employees inherit language and procedures without seeing the original evidential ambiguity. A contingent emergency response becomes policy, the policy becomes precedent, and the precedent becomes evidence that the original classification was sound.

Path dependence is not an argument against precaution. Institutions must sometimes act rapidly. It is an argument for marking provisional decisions, recording their evidential basis, scheduling review, and preventing the decision itself from being counted as new proof of the proposition that motivated it.

Otherwise an institution can climb a circular staircase: the public reaction justifies the institutional act, the act justifies stronger public reaction, and every turn feels like upward progress even though no new floor has been reached.

10.10 Selection against rational intervention

When interveners repeatedly receive the villain's vices, people who value precision face a choice: remain silent, leave the institution, narrow their work to safe topics, or accept chronic reputational cost. Over time, groups can select against the very dispositions required for self-correction.

The remaining membership need not become less intelligent. It can become more skilled at rationalizing the norm. Sophisticated participants learn which distinctions are forbidden, how to frame coalition-serving claims as methodological necessities, and how to punish dissent without naming the punishment. Intelligence improves the defense of the cascade rather than its accuracy.

This selection effect can create a misleading comparison. The institution points to the absence of internal dissent as proof that no serious dissent exists. But the absence may be an outcome of prior sanctions. Consensus is then both produced by the exclusion process and cited to justify it.

At a societal scale, repeated selection against interveners depletes a shared class of epistemic mediators: people trusted enough by more than one coalition to correct inconvenient claims. Without such mediators, every correction arrives as enemy speech, and even accurate institutions are interpreted through partisan ownership.

This is like selecting smoke detectors for silence. The organization may become calmer in the short run, but only because it has removed the devices capable of warning it.

11. Preserving Rationality in the Presence of a Villain

The remedy is not emotional numbness, endless delay, or fear of condemnation. It is a set of habits for the very moments when strong condemnation may be right and biased thinking is especially tempting.

11.1 Decompose the claim

Break a compound accusation into propositions. Instead of "the leader and everyone around the leader participated in the crime," ask separately:

1. What act occurred?
2. Who performed it?
3. What did the leader know?

4. When was it known?
5. What assistance, concealment, authorization, or negligence is alleged?
6. Which evidence supports each element?
7. What confidence is justified for each proposition?

Decomposition prevents a strong item from carrying weaker items invisibly. It also permits mixed conclusions. One can find concealment well supported, direct participation unproven, negligence likely, and a more sensational charge unsupported.

This is like unpacking a suitcase instead of weighing it while closed. A heavy suitcase tells us that something is inside. Only unpacking shows which items are actually there.

11.2 Use a claim-evidence ledger

For contested cases, a simple ledger can restore discipline:

Claim	Best supporting evidence	Best contrary evidence	Source status	Confidence	What would change confidence?
c1	item(s) bearing directly on c1	item(s) bearing against c1	primary, secondary, anonymous, disputed	low, moderate, high	specified observation

The ledger does not mechanize judgment. It prevents evidence for one proposition from migrating unnoticed to another. It also exposes whether the conclusion can be revised.

The point is not to turn moral judgment into bookkeeping. It is to stop a receipt for one item from being used as proof that every item in the cart was purchased.

11.3 Replace labels with descriptions, then restore labels if earned

Temporarily remove the loaded term and describe the conduct. Replace “traitor” with the alleged disclosure, agreement, or act of assistance. Replace “fascist” with the relevant position on political violence, elections, executive power, pluralism, or civil liberties. Replace “predator” with the specified behavior and evidence.

After the description is secure, ask whether the label adds accurate classification. This sequence preserves the possible value of moral vocabulary while preventing connotation from doing evidential work.

Describe the contents before choosing the label for the jar. If the contents meet the definition, use the strong label without apology. If they do not, a frightening label only hides what is actually inside.

11.4 Separate priors from proof

Prior conduct can increase or decrease plausibility. State that role openly: “Because the target has repeatedly done A, I assign a higher prior probability to related claim B.” Then ask what new evidence updates the prior.

This avoids two opposite errors. One is pretending character never matters. The other is treating character as proof. Bayesian language is not required, but the structure is useful: background information sets expectations; case-specific evidence changes confidence; neither should be disguised as the other.

A weather forecast raises the chance that the street will be wet tomorrow. Looking out the window tells us whether it is wet now. Prior conduct is the forecast; evidence about the present claim is the window.

11.5 Apply the source-switch and jersey-switch tests

Replace the target with an allied figure while holding the evidence fixed. Would the accusation still be accepted? Replace the accusing source with an opposing source. Would the report still be considered credible? Reverse the parties and ask whether the same procedural power would still seem legitimate.

These tests do not prove a conclusion false. They reveal whether identity is altering the rule. If the standard changes, the evaluator must identify a relevant difference or revise the judgment.

11.6 State both error costs

Before reaching a verdict, identify the harms of both false acceptance and false rejection. What happens if a true allegation is dismissed? What happens if a false allegation is accepted? Who bears each cost? Which precautionary actions are reversible? Which public statements are not?

This discipline prevents moral concern from being monopolized by one error. It is compatible with acting urgently when the expected harm of inaction is high. It simply requires that urgency be expressed through proportionate, reversible measures where possible.

Judgment often travels between two cliffs: failing to stop real harm and punishing someone for an unsupported claim. Seeing both cliffs does not force us to stand still. It helps us choose a safer path.

11.7 Distinguish protective action from final judgment

Institutions often must act before certainty. They can preserve evidence, limit access, appoint independent reviewers, warn potential victims, or impose temporary measures. The justification is risk management under uncertainty, not a fictional declaration that uncertainty has disappeared.

Transparent provisionality protects both safety and truth. It also makes later revision less humiliating. A community that can say “we acted because the risk was serious, and we are changing course because the evidence changed” is more resilient than one that must convert every precaution into a permanent verdict.

We already accept this distinction in ordinary life. A doctor may isolate a patient while waiting for a test. The isolation is a safety step, not a diagnosis manufactured by the locked door.

11.8 Make confidence visible

Use calibrated language. “Possible,” “plausible,” “more likely than not,” “well corroborated,” and “established under a stated standard” communicate different epistemic positions. Numerical probabilities are not always appropriate, but ordinal confidence is better than silent drift.

Calibration also separates moral clarity from factual certainty. One may say, “If this occurred, it would be atrocious; the present evidence makes occurrence plausible but not established.” The conditional condemnation need not be softened merely because the factual assessment remains open.

Confidence is a dimmer switch, not only an on-off button. “Possible,” “likely,” and “established” give the reader much more light than a vague glow of certainty.

11.9 Require a disconfirmation condition

Ask what evidence would reduce confidence. If no conceivable information would count, identify the belief as insulated. If only impossible evidence would count, the standard is functionally unfalsifiable.

The condition need not be a single decisive fact. Complex beliefs can be revised cumulatively. But the direction of possible revision must exist.

A belief with no door for contrary evidence is not a guarded conclusion. It is a locked room.

11.10 Protect evidential dissent

Communities need roles for internal critics who can challenge a favored accusation without being treated as allies of the accused. This protection should extend only to good-faith scrutiny, not harassment or sabotage, but the distinction must be based on conduct rather than conclusion.

The internal critic performs a public service. Adversaries will eventually test the claim. Discovering weaknesses early protects the credibility of sound accusations and prevents avoidable injustice.

Manufacturers test products before customers do. A movement should value internal checking for the same reason: an early repair is cheaper than a public collapse.

11.11 Redirect platform incentives toward accuracy

Pennycook and colleagues (2021) found that small prompts directing attention to accuracy improved the quality of news participants were willing to share, including in a field experiment on Twitter. The finding cautions against treating every false share as a settled false belief or an incurable partisan commitment. Attention is often captured by considerations other than truth.

Platforms and communities can make accuracy salient by requiring source links for grave claims, adding friction before resharing, separating allegation from verification in interface language, displaying corrections comparably to originals, and rewarding users who revise. No design eliminates motivated reasoning, but design helps determine which motive is cued at the moment of transmission.

A small speed bump before the “share” button can be useful. It does not ban travel; it gives attention a moment to catch up with emotion.

11.12 Build structured encounters

Deliberation is not a guaranteed cure; poorly structured group discussion can polarize. Yet carefully designed encounters can help. Fishkin, Siu, Diamond, and Bradburn (2021) reported that a national field experiment bringing more than 500 registered voters into informed, moderated deliberation produced substantial depolarizing changes in both issue positions and affective polarization.

The relevant features are not simply contact or civility. Participants received balanced materials, encountered actual people rather than villain abstractions, and deliberated under conditions that made reasons more rewarding than performance. Public institutions cannot place an entire society in one room, but they can reproduce elements of the epistemic design.

The lesson is closer to a well-run jury room than to a polite dinner. People need shared material, time, fair turns, and a reason to answer evidence instead of playing to an audience.

11.13 Audit the provenance of consensus

Before treating agreement as evidence, ask how many independent routes produced it. Did participants inspect primary material, rely on one report, infer from institutional action, or infer from one another’s confidence? Were contrary judgments visible? What sanctions attached to dissent? Does “everyone” refer to a representative field or to a self-selected room?

A consensus-provenance map can be as useful as a claim-evidence ledger. List the apparent sources, trace shared dependencies, separate expert method from prestige, and mark judgments formed before or after public controversy. The objective is not to discount all social knowledge. It is to avoid counting echoes as witnesses.

Draw the consensus as a family tree. If twenty branches all lead back to one trunk, the picture shows spread, not twenty independent roots.

Institutions should apply the same discipline to public pressure. A large volume of messages may reflect broad independent concern, coordinated mobilization, one viral item, or automated repetition. Each can matter politically, but they do not carry the same epistemic meaning.

11.14 Separate an epistemic role from its immediate beneficiary

Evaluate an intervention in two stages. First ask whether its reasoning is sound. Only then ask whether the intervener is acting strategically or has a conflict that changes the weight of the contribution. Do not derive the quality of the reason from the villainized identity of the person who benefits.

A useful formulation is: “This correction favors *t* with respect to *c*; it does not establish approval of *t*, alignment with *t*, or similarity to *t*.” Making the limited implication explicit interrupts vice transfer. The same principle should govern formal roles. A favorable evidential finding is not a character certificate, and a defense of procedure is not a defense of the conduct under review.

An umpire can make one correct call for a player who cheated earlier in the game. The call does not wash the player clean, and the earlier cheating does not make the call wrong.

Where complicity is alleged, decompose it like any other claim. Specify the assistance, knowledge, intent, causal contribution, violated duty, and evidence for each. This protects the concept of complicity from dilution and permits real collaboration with wrongdoing to be condemned more credibly.

11.15 Lower the threshold for honest dissent

Because a cascade is partly a coordination problem, private permission is insufficient. Organizations need visible channels through which qualified disagreement can become common knowledge. Anonymous pre-decision polling, recorded minority reports, rotating red-team roles, independent ombuds functions, and confidential review can reveal whether public unanimity masks dispersed doubt.

The design must also prevent retaliation by association. Evaluations of the dissenter should cite conduct: fabrication, harassment, conflicts, refusal to engage evidence, or breach of duty. “This objection aids the accused” should never be allowed to stand as a character finding.

One visible dissenter can reduce the perceived isolation of others. The goal is not to reward contrarianism or require artificial balance. It is to make the cost of expressing evidence track the quality of the evidence rather than the identity of its beneficiary.

One person raising a hand can show ten silent people that they were not alone. The value lies less in automatic disagreement than in making hidden doubt visible.

11.16 Create exit ramps for correction

People resist correction when revision requires repudiating their entire moral identity. Communities should make narrower repair possible. A speaker should be able to say, “The target remains culpable for A; I was wrong about B; my correction of B does not weaken the case for A.” Proposition ownership and decomposition create this exit ramp.

Corrections should match the reach and clarity of the original claim, identify the exact change, and avoid preserving a false detail through vague language such as “questions have been raised.” Institutions should timestamp provisional decisions, schedule reassessment, and keep the original evidential basis visible so that later action is not mistaken for independent proof.

Leaders can shift injunctive norms by rewarding accurate revision publicly. When a respected member corrects an ally’s exaggeration without losing group standing, observers learn that rational intervention is compatible with loyalty. This is how the irrationality ratchet can be reversed: not only by supplying new facts, but by changing what the group treats as honorable.

An exit ramp makes turning around less humiliating than crashing through a barrier. Communities need language and customs that let people repair one claim without pretending their whole moral concern was false.

12. Objections and Replies

Objection 1: Some villains really are villains.

Yes. The paper does not argue that moral condemnation is usually mistaken or that character never matters. It argues that a true judgment about a person’s general conduct does not prove every particular charge. A target’s real badness makes care more important because nearby false claims will be unusually easy to believe.

The point can be stated without the word “villain”: no amount of proven wrongdoing guarantees that a different allegation is true. A thief may be a reasonable suspect when money disappears. The earlier theft still does not tell us who took this money.

Objection 2: Public discourse is not a court.

Correct. Different settings require different standards, procedures, and consequences. A private citizen need not write a judicial opinion before expressing suspicion. Journalism, friendship, voting, organizational safety, and criminal adjudication properly operate differently.

But “not a court” does not mean “not answerable to evidence.” A factual assertion remains capable of being true or false outside a courtroom. The appropriate public standard may be lower than proof beyond a reasonable doubt, but it is not zero. In fact, the absence of courtroom safeguards increases the speaker’s responsibility to distinguish allegation from finding.

A neighborhood basketball game is not the national league, but the score still cannot be invented. Informal discussion needs fewer formal steps than a court, not freedom from truth.

Objection 3: Cumulative evidence can establish a case without a single decisive item.

Certainly. Many sound conclusions rest on convergent circumstantial evidence. The criticism of the evidential heap is not a demand for a smoking gun. It is a demand that the pieces bear on the same proposition, that their dependence be assessed, that contrary evidence be considered, and that the conclusion not be silently changed.

Ten modestly independent items can support a conclusion strongly. Ten repetitions of one allegation do not. Ten facts showing that a person is generally corrupt do not automatically establish one specified transaction. The difference is inferential structure, not drama.

Many strands can make a strong rope when each adds strength. Ten loops made from the same frayed strand only look thicker.

Objection 4: Demanding precision benefits oppressors.

Precision can be weaponized. Powerful actors may demand unattainable evidence, exploit institutional delay, hide records, and portray every discrepancy as disproof. A responsible account must recognize evidential obstruction.

But imprecision also benefits oppressors. It allows them to merge strong accusations with weak ones and discredit the whole set by exposing the weakest. It enables retaliation through fabricated claims. It trains supporters to accept accusation by identity, a rule that power can use more effectively than dissidents can.

Precision is not passivity. One can specify that evidence was destroyed, access was unequal, testimony converges despite missing records, or precaution is warranted under uncertainty. The demand is not perfect information; it is honest representation of the information available.

A foggy photograph may still show enough to act, and a missing record may itself need explanation. Precision asks us to say, “The picture is foggy, but these features are visible,” rather than pretending the picture is sharp.

Objection 5: Victims should be believed.

Victims should be heard, protected from retaliation, treated with dignity, and provided competent avenues of investigation. Their testimony is evidence. None of this makes every testimony infallible or every interpretation exact.

The phrase “believe victims” is often best understood as a corrective to reflexive dismissal rather than a literal rule of automatic factual acceptance. Interpreted that way, it is compatible with serious inquiry. Interpreted as a rule that self-identification determines truth, it cannot manage conflicting allegations, mistaken identification, memory limits, or strategic fabrication.

Objection 6: Both-sides framing conceals asymmetries.

A symmetrical psychological capacity does not imply symmetrical behavior, power, media structure, or harm. One faction can produce more misinformation in a particular ecosystem. One state can possess greater coercive power. One movement can violate democratic rules more severely. Empirical claims about these asymmetries should be investigated rather than neutralized.

The paper’s symmetry requirement concerns standards. If the same evidence would be rejected against an ally and accepted against an opponent, the asymmetry requires justification. Applying one procedure to unequal facts will often produce unequal conclusions. That is not false equivalence; it is what permits real differences to emerge credibly.

Using the same ruler does not mean everyone has the same height. It is what allows real differences in height to be measured.

Objection 7: Anger and disgust are morally informative.

They can be. Emotion draws attention to harm, communicates commitment, motivates collective action, and registers values that detached calculation may miss. The argument is not an attempt to replace citizens with affectless machines.

The problem arises when intensity of emotion is treated as evidence of the event, when emotional display becomes a loyalty requirement, or when disgust reduces the target's eligibility for justice. Emotion should help identify what matters. It should not decide what happened without evidential mediation.

Smoke can alert us to look for fire. The strength of our alarm does not tell us whether the smoke came from a burning house or a burnt pan.

Objection 8: Refusing maximal language normalizes evil.

Sometimes ordinary language understates exceptional danger. Historical accuracy may require words such as "fascist," "terrorist," "apartheid," "torture," or "genocide." Refusing an earned term can itself be a form of evasion.

The connotative audit does not require euphemism. It asks the speaker to support the classification and use the same criteria across targets. A term that survives that audit becomes more powerful, not less. It no longer functions as thrown dung but as a defensible description.

A warning sign is strongest when it names the actual danger. Writing every possible hazard on every sign does not make the public safer.

Objection 9: The target's pattern makes further misconduct overwhelmingly likely.

Perhaps. Then the pattern should be specified, its base rate considered, and its relevance to the new allegation explained. "Overwhelmingly likely" is already a probabilistic claim and can be assessed.

Even a high probability is not identity. If nine of ten similar actions were corrupt, one should investigate the tenth aggressively. One should not falsify the record by saying the tenth has been established before examining it.

A deck loaded with red cards makes red likely on the next draw. It does not make the unseen card red before anyone turns it over.

Objection 10: Procedural universalism protects the powerful more than the weak.

Formal rules can operate unequally when access to counsel, records, media, and institutional attention is unequal. Universal procedure without material access can preserve injustice. The answer is to improve access, independence, transparency, witness protection, investigative competence, and accountability for obstruction.

Target-indexed standards do not solve structural inequality. They replace a contest over evidence with a contest over who can successfully assign villain status. Powerful actors are often best positioned to win that contest.

Fair rules without fair access are like giving everyone the same race instructions while some begin far behind the starting line. The answer is to repair access as well as keep the finish line fixed.

Objection 11: A large consensus is itself strong evidence.

Often it is. Convergent judgments from people with relevant competence, adequate access, sound methods, and meaningful independence can be among the best evidence available. The problem is not inference from consensus but inference from an unanalyzed appearance of consensus.

Agreement generated by one source, a shared incentive, public conformity, or fear of vice transfer carries less information than superficially similar agreement generated independently. The rational question is not

simply how many agree. It is how the agreement was produced. A consensus-provenance audit can strengthen genuine expert convergence by distinguishing it from social repetition.

Twenty thermometers provide strong evidence only if they are actually taking readings. Twenty screens displaying the output of one broken thermometer do not.

Objection 12: The intervener really does help the villain.

Sometimes a correction, delay, or procedural ruling benefits a guilty and dangerous person. Consequences matter. The mistake is treating benefit as a complete moral classification of the act or actor. Procedures produce local benefits for bad people because they are designed to reduce errors across cases.

One should ask whether the intervention is accurate, proportionate, consistently applied, and attentive to risk. If a less harmful way to correct the record exists, use it. If immediate precaution is justified, take it. But an accurate correction does not become false because of its beneficiary, and the beneficiary's vice does not migrate to the corrector.

Repairing the brakes on a criminal's car does not make the mechanic a criminal. The work may help a bad person, yet the mechanic's honesty and the brake test remain separate questions.

Objection 13: The intervener's tone, persistence, or strange priorities reveal the hidden vice.

Tone and pattern can supply evidence about motive. A person who questions only accusations against one faction, ignores decisive answers, harasses complainants, or obsessively minimizes one kind of harm may deserve scrutiny. The paper does not require pretending that behavior has no diagnostic value.

But diagnosis must remain comparative and defeasible. Was the persistence caused by repeated nonanswers? Is the standard applied elsewhere? Does the person acknowledge strong evidence? Is the alleged vice inferred from observable conduct or merely restated from the disagreement? An irritating critic can be right, and a courteous consensus can be wrong. Character assessment cannot substitute for answering the claim.

A smoke alarm can be shrill, repetitive, and badly timed while still detecting smoke. Tone may affect how we respond to the alarm; it does not settle whether smoke is present.

Objection 14: Harsh conformity is necessary when a community faces a genuine threat.

Communities confronting infiltration, violence, abuse, or organized deception cannot treat every intervention as innocent. Coordination, secrecy, and rapid action may be necessary. A norm against irresponsible doubt can protect vulnerable people from strategic obstruction.

Yet threat increases the value of reliable discrimination. A community that cannot distinguish an infiltrator from an internal corrector can be manipulated by planted accusations, purge competent members, and lose access to unfavorable information. Emergency action may use lower provisional thresholds, but it should remain scoped, reviewable, and explicit about uncertainty. The more coercive the consequence, the less safe it is to make membership performance substitute for evidence.

An immune system that attacks every unusual cell can harm the body it protects. A threatened community also needs a way to distinguish danger from its own error-correction system.

13. An Empirical Research Program

The villain-exception idea can be tested rather than left as an interesting description.

13.1 Basic factorial design

Participants could evaluate a novel accusation against a public figure in a 3 x 2 x 2 design:

- target moral history: admirable, neutral, or documented serious misconduct;
- partisan identity: aligned or opposed to the participant;
- evidence quality: weak or strong.

The main dependent variables would include confidence in the allegation, amount of additional evidence requested, willingness to share, support for investigation, support for punishment, perceived immorality of a person requesting more evidence, and recall of allegation status. Belief, precaution, and punishment should be measured separately so that a lower threshold for protective action is not mistaken for a lower threshold for factual belief.

The villain-exception prediction is not merely that prior misconduct raises confidence. Some increase may be rational. It predicts an interaction: documented misconduct will reduce sensitivity to the difference between weak and strong case-specific evidence, especially for politically opposed targets, and will increase moral condemnation of evidential dissent. Stimuli should be pretested to verify that the weak and strong evidence conditions are distinguishable without producing floor or ceiling effects.

For example, two made-up mayors could face the same new charge. One has a clean record; the other has a history of serious misconduct. If participants treat the same weak evidence almost like strong evidence only for the second mayor, that would reveal more than ordinary use of background information.

13.2 Claim-specificity manipulation

A second design could vary whether the prior misconduct is closely related or merely emotionally adjacent to the new allegation. If participants rationally use character evidence, closely related priors should produce more updating than unrelated but equally disgusting misconduct. If disgust supplies a general villain discount, emotionally intense unrelated wrongdoing may inflate confidence beyond its evidential relevance.

13.3 Label-scatter manipulation

Participants could receive identical behavioral descriptions paired with:

1. one precisely defined severe label;
2. a list of multiple severe labels;
3. neutral descriptive language;
4. a list containing one accurate and several inaccurate labels.

Researchers could measure perceived evidence strength, memory for the underlying facts, willingness to punish, and confidence that each label was individually established. The connotative-scattershot hypothesis predicts that a label cluster will intensify global condemnation even when participants cannot match evidence to the individual terms.

13.4 Closed-system detection

Participants could be shown several possible target responses: denial, silence, detailed rebuttal, partial concession, or disclosure of records. If all responses increase perceived guilt for villain targets but not neutral targets, that would support the hypothesis that the accusation has become insulated from disconfirmation.

13.5 Social-validation and intensity-matching experiment

Participants could enter simulated discussion networks in which earlier comments express low, medium, or high certainty about an accusation supported by identical weak evidence. A second manipulation could vary social rewards: no feedback, approval for calibrated comments, or approval for maximal condemnation.

Researchers could measure the participant's private confidence, public wording, label severity, evidence requested, and willingness to sanction a dissenter. The irrationality-ratchet hypothesis predicts that visible certainty and reward norms will widen the gap between private and public confidence, reduce discrimination between weak and strong evidence, and increase the intensity participants believe necessary to appear morally adequate.

In plain terms, three rooms could see the same thin news item. One room rewards careful comments, one rewards absolute condemnation, and one gives no feedback. The question is whether the same evidence begins to sound stronger simply because the room applauds stronger language.

Sequential designs could also test path dependence. Networks would receive the same total evidence but encounter early comments in different orders. If early maximalism produces later categorical consensus, the result would distinguish sociological momentum from mere aggregation of stable individual judgments.

13.6 Vice-transfer and rational-intervener experiment

Participants could read an accusation followed by an intervention that either requests a source, corrects a factual detail, defends a procedural rule, strategically delays an inquiry, or actively conceals evidence. The intervener could be an ingroup member, outgroup member, or politically neutral professional. The target's attributed vice could be established by evidence, merely alleged, or absent.

Dependent variables would include evaluation of the intervention's reasoning, perceived allegiance, attribution of the target's vice to the intervener, support for social punishment, and willingness to consider the correction. Vice transfer predicts an inferential progression from favorable effect to allegiance to shared vice even in the source-request and factual-correction conditions. A crucial test is whether an unestablished vice is assigned to the intervener more confidently than it was assigned to the target.

A very simple intervention might correct a date or ask for the original document. Researchers could then see whether participants judge that small act on its merits or treat it as a moral handshake with the villain.

The design should include genuine obstruction conditions so that resistance to vice transfer is not confused with blindness to complicity. A successful measure must discriminate principled scrutiny from concealment rather than simply making participants more charitable toward every critic.

13.7 Pluralistic-ignorance and dissent-threshold study

Researchers could privately measure group members' confidence in a controversial claim, then vary what participants are told about others' private confidence and public expression. Some groups would receive accurate norm information, others inflated public certainty, and others an anonymous distribution of actual private judgments.

Outcomes would include subsequent public expression, perceived deviance of uncertainty, willingness to join a minority report, and estimates of group consensus. If counterfeit consensus sustains the villain exception, revealing the true distribution of private confidence should reduce performative certainty without necessarily reducing sensitivity to strong evidence.

Suppose a room sounds 90 percent certain in public but its anonymous poll shows an average confidence of only 55 percent. Revealing the poll would test whether much of the certainty was resting on a mistaken picture of everyone else.

A field version could introduce confidential polling or protected minority reports in organizations before high-stakes statements. The important outcome would be whether these devices improve later calibration and correction, not whether they mechanically produce moderation.

13.8 Symmetry and cross-national scope

Studies should pre-register political reversals and sample multiple ideological groups. Stimuli should include leaders and movements from the left, right, religious, secular, nationalist, anti-nationalist, establishment, and anti-establishment contexts. The aim would not be to declare all groups equivalent but to identify when target identity changes evidential sensitivity.

Cross-national work is important because party systems, media ecologies, legal cultures, and histories of repression vary. “Villain” may be attached to different categories, while the threshold mechanism remains similar.

13.9 Interventions

Promising interventions include:

- a source-switch prompt;
- a request to list claim elements;
- an explicit reminder of both error costs;
- confidence calibration;
- a requirement to identify disconfirming evidence;
- replacement of labels with behavioral descriptions;
- an accuracy prompt before sharing;
- exposure to the opposing group’s actual, rather than imagined, attitudes;
- structured cross-partisan deliberation.

Outcomes should include not only reduced belief. A good intervention should increase discrimination between strong and weak evidence. General skepticism is not the goal. The rational mind should become more responsive to evidence, not merely less willing to conclude.

Studies should measure public and private judgments separately. A treatment may leave belief unchanged while reducing the pressure to perform certainty, or it may improve private calibration without changing public expression because sanctions remain. Conflating these outcomes would miss the sociological component of the theory.

13.10 What would count against the theory

A useful framework must risk being wrong. Evidence would count against a broad villain-exception claim if documented misconduct raised reasonable prior suspicion but left participants equally sensitive to weak and strong case-specific evidence, equally willing to revise, and no more hostile toward a good-faith critic. A result in which villainized targets receive harsher moral judgments but unchanged evidential thresholds would support ordinary condemnation, not the proposed exception.

Studies should therefore preregister the predicted interactions, validate the evidence-strength and villain-status manipulations, report null results, and replicate across topics and political directions. The aim is not to find any sign that disliked targets are judged more harshly. It is to isolate the narrower failure in which target identity changes the handling of evidence or the treatment of correction.

14. Conclusion

The villain is not only a person in a story. The villain can become a thinking shortcut. Once a target seems bad enough, ordinary rules begin to look optional. A rumor is treated as a fact, possibility becomes certainty, association becomes participation, and a pile of ugly labels stands in for a careful description. Punishment is demanded before the charge is made clear. No path back is left open.

The psychology behind this is easy to recognize. We cross-examine claims we dislike and give friendly interviews to claims we want to believe. Moral anger turns the volume up until small differences are hard to hear. The belief that someone is evil turns a person into a stencil through which every new event is traced. Party and group identity then reward conclusions that help our side.

Groups keep the pattern alive. A crowd works like a stadium wave: each person moves partly because nearby people moved. Repeated claims start to look like independent witnesses even when they are only echoes. People hide their doubts because everyone else sounds certain, while everyone else may be doing the same. The public temperature rises as members compete to show that they are fully committed. By the time anyone checks the original evidence, the claim may already be tied to friendships, jobs, sanctions, news archives, and institutional decisions.

The rational intervener threatens this arrangement by asking a plain question: What does the evidence support? The intervener need not admire or excuse the villain. Yet the question may help the villain in one limited dispute, so help is turned into allegiance and allegiance into shared vice. The referee is accused of joining the dirty team because one correct call favored it. If the villain's vice was only rumored, the rumor is copied onto a second person. If the vice was real, it is still transferred without evidence.

Nothing in this paper shows that serious accusations are generally false. The point is almost the reverse. Real wrongdoing deserves a clear record. A villain's history can be a smoke alarm that justifies quick attention. It is not a camera recording every later event. Genuine victims and sound accusations are harmed when they must compete with exaggerations dressed in the same language.

The public cost grows with each repetition. Precision starts to look disloyal. Correction looks like aid to the enemy. Institutions are pushed to announce conclusions before finishing their work. Their actions are then cited as proof that the original claim was true. Each side keeps the other's excesses as permission slips for its own. Citizens eventually conclude that standards are team uniforms: worn proudly when useful and changed when the other side takes the field.

The alternative is not timidity. It is moral seriousness with better workmanship. Condemn what the evidence establishes. Investigate what it makes plausible. Protect people when the risk justifies temporary action. Say what remains unknown. Trace a crowd's many claims back to their actual roots. Describe conduct before reaching for the largest label. Give internal critics a safe way to speak. Correct an error as publicly as it was spread. Keep the rule after switching jerseys.

Anyone can defend accuracy when it helps a friend. The harder test comes when accuracy helps someone we despise. A fair referee still calls the play correctly. A reliable alarm does not change its threshold according to who owns the house. Standards that work only in the absence of villains are not standards. A rational community must oppose real wrongdoing, reject invented wrongdoing, and protect the people willing to tell the difference.

References

- Bandura, A. (1999). Moral disengagement in the perpetration of inhumanities. *Personality and Social Psychology Review*, 3(3), 193-209. https://doi.org/10.1207/s15327957pspr0303_3
- Bikhchandani, S., Hirshleifer, D., & Welch, I. (1992). A theory of fads, fashion, custom, and cultural change as informational cascades. *Journal of Political Economy*, 100(5), 992-1026. <https://doi.org/10.1086/261849>
- Brady, W. J., McLoughlin, K., Doan, T. N., & Crockett, M. J. (2021). How social learning amplifies moral outrage expression in online social networks. *Science Advances*, 7(33), Article eabe5641. <https://doi.org/10.1126/sciadv.abe5641>
- Brady, W. J., McLoughlin, K. L., Torres, M. P., Luo, K. F., Gendron, M., & Crockett, M. J. (2023). Overperception of moral outrage in online social networks inflates beliefs about intergroup hostility. *Nature Human Behaviour*, 7, 917-927. <https://doi.org/10.1038/s41562-023-01582-0>
- Brady, W. J., Wills, J. A., Jost, J. T., Tucker, J. A., & Van Bavel, J. J. (2017). Emotion shapes the diffusion of moralized content in social networks. *Proceedings of the National Academy of Sciences*, 114(28), 7313-7318. <https://doi.org/10.1073/pnas.1618923114>
- Cialdini, R. B., Reno, R. R., & Kallgren, C. A. (1990). A focus theory of normative conduct: Recycling the concept of norms to reduce littering in public places. *Journal of Personality and Social Psychology*, 58(6), 1015-1026. <https://doi.org/10.1037/0022-3514.58.6.1015>
- Deutsch, M., & Gerard, H. B. (1955). A study of normative and informational social influences upon individual judgment. *Journal of Abnormal and Social Psychology*, 51(3), 629-636. <https://doi.org/10.1037/h0046408>
- Ditto, P. H., & Lopez, D. F. (1992). Motivated skepticism: Use of differential decision criteria for preferred and nonpreferred conclusions. *Journal of Personality and Social Psychology*, 63(4), 568-584. <https://doi.org/10.1037/0022-3514.63.4.568>
- Ditto, P. H., Liu, B. S., Clark, C. J., Wojcik, S. P., Chen, E. E., Grady, R. H., Celniker, J. B., & Zinger, J. F. (2019). At least bias is bipartisan: A meta-analytic comparison of partisan bias in liberals and conservatives. *Perspectives on Psychological Science*, 14(2), 273-291. <https://doi.org/10.1177/1745691617746796>
- Fazio, L. K., Brashier, N. M., Payne, B. K., & Marsh, E. J. (2015). Knowledge does not protect against illusory truth. *Journal of Experimental Psychology: General*, 144(5), 993-1002. <https://doi.org/10.1037/xge0000098>
- Fishkin, J., Siu, A., Diamond, L., & Bradburn, N. (2021). Is deliberation an antidote to extreme partisan polarization? Reflections on "America in One Room." *American Political Science Review*, 115(4), 1464-1481. <https://doi.org/10.1017/S0003055421000642>
- Graham, M. H., & Svolik, M. W. (2020). Democracy in America? Partisanship, polarization, and the robustness of support for democracy in the United States. *American Political Science Review*, 114(2), 392-409. <https://doi.org/10.1017/S0003055420000052>
- Granovetter, M. (1978). Threshold models of collective behavior. *American Journal of Sociology*, 83(6), 1420-1443. <https://doi.org/10.1086/226707>
- Gray, K., & Wegner, D. M. (2009). Moral typecasting: Divergent perceptions of moral agents and moral patients. *Journal of Personality and Social Psychology*, 96(3), 505-520. <https://doi.org/10.1037/a0013748>
- Iyengar, S., & Westwood, S. J. (2015). Fear and loathing across party lines: New evidence on group polarization. *American Journal of Political Science*, 59(3), 690-707. <https://doi.org/10.1111/ajps.12152>
- Kteily, N., Bruneau, E., Waytz, A., & Cotterill, S. (2015). The ascent of man: Theoretical and empirical evidence for blatant dehumanization. *Journal of Personality and Social Psychology*, 109(5), 901-931. <https://doi.org/10.1037/pspp0000048>
- Kunda, Z. (1990). The case for motivated reasoning. *Psychological Bulletin*, 108(3), 480-498. <https://doi.org/10.1037/0033-2909.108.3.480>
- Marques, J. M., Yzerbyt, V. Y., & Leyens, J.-P. (1988). The "black sheep effect": Extremity of judgments towards ingroup members as a function of group identification. *European Journal of Social Psychology*, 18(1), 1-16. <https://doi.org/10.1002/ejsp.2420180102>
- Monin, B., Sawyer, P. J., & Marquez, M. J. (2008). The rejection of moral rebels: Resenting those who do the right thing. *Journal of Personality and Social Psychology*, 95(1), 76-93. <https://doi.org/10.1037/0022-3514.95.1.76>

- Moore-Berg, S. L., Ankori-Karlinsky, L. O., Hameiri, B., & Bruneau, E. G. (2020). Exaggerated meta-perceptions predict intergroup hostility between American political partisans. *Proceedings of the National Academy of Sciences*, 117(26), 14864-14872. <https://doi.org/10.1073/pnas.2001263117>
- Moscovici, S., & Zavalloni, M. (1969). The group as a polarizer of attitudes. *Journal of Personality and Social Psychology*, 12(2), 125-135. <https://doi.org/10.1037/h0027568>
- Pennycook, G., Epstein, Z., Mosleh, M., Arechar, A. A., Eckles, D., & Rand, D. G. (2021). Shifting attention to accuracy can reduce misinformation online. *Nature*, 592, 590-595. <https://doi.org/10.1038/s41586-021-03344-2>
- Prentice, D. A., & Miller, D. T. (1993). Pluralistic ignorance and alcohol use on campus: Some consequences of misperceiving the social norm. *Journal of Personality and Social Psychology*, 64(2), 243-256. <https://doi.org/10.1037/0022-3514.64.2.243>
- Pryor, J. B., Reeder, G. D., & Monroe, A. E. (2012). The infection of bad company: Stigma by association. *Journal of Personality and Social Psychology*, 102(2), 224-241. <https://doi.org/10.1037/a0026270>
- Salerno, J. M., & Peter-Hagene, L. C. (2013). The interactive effect of anger and disgust on moral outrage and judgments. *Psychological Science*, 24(10), 2069-2078. <https://doi.org/10.1177/0956797613486988>
- Taber, C. S., & Lodge, M. (2006). Motivated skepticism in the evaluation of political beliefs. *American Journal of Political Science*, 50(3), 755-769. <https://doi.org/10.1111/j.1540-5907.2006.00214.x>
- Van Bavel, J. J., Packer, D. J., Haas, I. J., & Cunningham, W. A. (2012). The importance of moral construal: Moral versus non-moral construal elicits faster, more extreme, universal evaluations of the same actions. *PLOS ONE*, 7(11), Article e48693. <https://doi.org/10.1371/journal.pone.0048693>
- Webster, R. J., & Saucier, D. A. (2013). Angels and demons are among us: Assessing individual differences in belief in pure evil and belief in pure good. *Personality and Social Psychology Bulletin*, 39(11), 1455-1470. <https://doi.org/10.1177/0146167213496282>

Appendix A. A Practical Accusation Audit

This appendix turns the paper’s argument into a working procedure. It is not a device for delaying every decision or giving powerful people endless chances to obstruct. It is a way to keep belief, precaution, investigation, and punishment from collapsing into one emotional verdict.

A.1 Put the claim on paper

Before debating a heated accusation, complete the following ledger. If a field cannot yet be filled, write “unknown” rather than letting the villain narrative fill it silently.

Field	What to record	Warning sign
Exact proposition	Who is alleged to have done what, when, and with what state of mind?	The claim changes when challenged
Claim status	Report, allegation, corroborated claim, formal finding, or final adjudication	These stages are treated as synonyms
Best supporting evidence	The strongest item bearing directly on the proposition	A large archive is named but no item is located
Best contrary evidence	The strongest item that lowers confidence or narrows the claim	Contrary material is dismissed because of who benefits
Source path	Primary source and every important intermediary	Many outlets trace back to one allegation
Confidence	Possible, plausible, likely, well corroborated, or established under a stated standard	Certainty is expressed without a matching threshold
Decision requested	Listen, preserve records, investigate, restrict access, publish, punish, or adjudicate	A severe sanction appears before a clear charge
Revision condition	What observation would raise, lower, or narrow confidence?	No possible evidence could change the judgment

The point is not paperwork for its own sake. A grocery receipt separates items that arrived in one bag. The ledger does the same for accusations that arrived in one emotional package.

A.2 Separate belief, precaution, and punishment

Ask three different questions:

1. What should a person currently believe?
2. What temporary protective action is justified under uncertainty?
3. What punishment or durable public judgment is justified?

The answers need not match. A school may separate two people while investigating without declaring either person guilty. A newsroom may investigate a plausible report without publishing it as established fact. A court may preserve evidence without deciding the case. Confusing these stages allows a sensible precaution to return later wearing a false badge of proof.

A.3 Run the reversal tests

Use four quick switches:

- **Jersey switch:** Would the evidential standard remain the same if the target belonged to the favored faction?
- **Source switch:** Would the source still seem reliable if the same method produced an unwelcome conclusion?
- **Beneficiary switch:** Would the correction still be judged on its merits if it helped a friend rather than a villain?
- **Sanction switch:** Would the procedure still seem fair if an opponent controlled it?

A result need not be identical after every switch. Groups and individuals can have different histories. The test asks whether a relevant difference explains the change or whether identity quietly changed the rule.

A.4 Audit the crowd

When “everyone knows” appears, ask:

1. How many independent routes produced the judgment?
2. Did participants inspect primary material or copy one another?
3. Were contrary judgments visible?
4. What social rewards attached to certainty?
5. What sanctions attached to doubt?
6. Does the public display match what members say privately?

One hundred retweets can represent broad independent concern, coordinated mobilization, one viral item, or automated repetition. Volume is information about reach. It is not, by itself, information about truth.

A.5 Leave an exit ramp

A correction should identify the original wording, the repaired wording, the reason for the change, and the reach of both versions. Institutions should preserve the original evidential basis so that revision does not look like unexplained surrender. A group that gives members no honorable way to correct an overstatement teaches them to defend it.

The minimum exit ramp is a sentence such as: “The target remains culpable for A; the evidence does not establish B; our earlier wording joined them too quickly.” This preserves moral judgment where warranted while reopening the door to accuracy.

Appendix B. Script-to-Safeguard Crosswalk

The field-guide scripts can be recognized by the substitution they perform. The first corrective question below is not a complete rebuttal. It is a way to return the discussion to the missing step.

B.1 Identity, evidence, and consensus

No.	Script family	Hidden substitution	First corrective question
1	Loyalty oath	Factional allegiance replaces assessment of the proposition	Which exact part of the charge must be repeated, and what establishes it?
2	Decency test	The critic's character replaces the missing proof	What evidence would answer the question regardless of who asked it?
3	Procedure as allegiance	A rule's beneficiary is treated as the rule's purpose	Is the procedure fair, workable, and applied across targets?
4	Beneficiary fallacy	Who benefits from a correction determines whether it is true	Is the correction accurate, and what independent evidence bears on it?
5	Counterfeit consensus	Visible repetition replaces independent agreement	How many independent sources or methods produced the conclusion?
6	Humiliation	An insult replaces the requested citation	What source, page, and proposition answer the request?
7	Self-sealing association	Denial is converted into confirmation	What possible response or evidence would lower confidence?
8	Compassion trap	Care for harmed people is equated with immediate belief	What protection is justified now, and what remains to be established?

B.2 Closure, character, and sanction

No.	Script family	Hidden substitution	First corrective question
9	Condemnation without a charge	Global badness replaces identification of the act	Which offense is alleged, and which evidence supports each element?
10	Prestige laundering	Status replaces relevant expertise, access, and method	What does this authority know, and how was it learned?
11	Complexity as confession	The form of the answer replaces assessment of its content	Which premise, source, or inference in the explanation is wrong?
12	Technical-language suspicion	The use of distinctions is treated as evasion	Can the distinction be stated plainly, and does it change the conclusion?
13	Diagnostic trap	Every response is made evidence of a corrupt identity	What response would count as good faith or disconfirmation?
14	Joke ratchet	The claim becomes literal for effect and figurative for responsibility	What proposition was the audience meant to take away?
15	Punishment first	The requested sanction supplies gravity to an unspecified charge	What act, rule, authority, and proportionate remedy are being proposed?
16	Expulsion as victory	Control of the room replaces an answer to the argument	What claim remains unanswered after the speaker is removed?

Several scripts can operate in one sentence. “Everyone knows what people like you are” combines counterfeit consensus, type assignment, and a closed diagnosis. The crosswalk is therefore a set of lenses, not sixteen boxes into which every exchange must fit.

Appendix C. Reversible Vignette and Coding Template

This appendix provides a reusable scaffold for testing the paper’s narrower empirical claims. All names, offices, factions, and allegations should be fictional. The design should be reversed across political directions so that “villain” is a manipulated role rather than a disguised label for one real coalition.

C.1 Vignette scaffold

Target and history: [Name] is a [role]. Independent records describe the target’s prior history as [admirable, neutral, or documented serious misconduct].

New allegation: A new source claims that the target [precisely specified act].

Evidence condition: The report is supported by [weak or strong claim-specific evidence, pretested for perceived strength].

Alignment condition: The target is described as [aligned with or opposed to] the participant’s relevant political or social identity.

Intervention condition, when used: A second speaker [requests the original source, corrects a factual detail, defends a procedural rule, strategically delays inquiry, or conceals evidence].

The misconduct history should raise reasonable prior suspicion without answering the new allegation. The weak and strong evidence conditions should concern the same proposition and differ in diagnostic value, not merely in length or emotional detail.

C.2 Core manipulation matrix

Dimension	Levels	Purpose
Target history	Admirable, neutral, documented serious misconduct	Distinguishes ordinary prior updating from a villain exception
Partisan or group alignment	Aligned, opposed	Tests directional symmetry and motivated evaluation
Claim-specific evidence	Weak, strong	Measures sensitivity to evidential quality
Intervener conduct, optional extension	Principled correction, neutral question, genuine obstruction	Separates vice transfer from warranted scrutiny of complicity
Public norm, optional extension	Accuracy rewarded, maximal condemnation rewarded, no feedback	Tests social validation and intensity matching

C.3 Outcome families

Do not compress all responses into one approval score.

Outcome family	Example measures	Why it must remain separate
Factual belief	Confidence in the allegation; memory for allegation status	Belief can remain uncertain while action is justified
Precaution	Support for preserving records, limiting access, or opening an investigation	Precaution responds to risk and both error costs
Punishment	Support for firing, removal, public condemnation, or legal sanction	Durable sanctions normally require a stronger basis
Revision	Evidence requested; willingness to lower confidence; recognition of correction	The theory predicts resistance to updating, not merely dislike
Intervener judgment	Perceived allegiance, vice, intelligence, and good faith	Directly tests vice transfer and the second villain
Social expression	Public wording, sharing, private confidence, and perceived group norm	Separates private judgment from performative certainty

C.4 Manipulation and quality checks

- Pretest whether the evidence conditions are distinguishable without floor or ceiling effects.
- Verify that the misconduct manipulation changes target evaluation without silently adding evidence about the new act.
- Ask participants to identify whether the item was a report, allegation, corroborated claim, or finding.
- Measure perceived source independence rather than counting visible repetitions.
- Record political identity before exposure when feasible, and preregister exclusion and coding rules.
- Include attention and comprehension checks that do not reveal the favored hypothesis.
- Report null interactions and results that support ordinary prior updating without supporting a villain exception.

C.5 Interpretation rule

Harsher judgment of a documented wrongdoer does not by itself support the theory. Stronger evidence should increase belief, and relevant history may rationally affect a prior. The clearest support would be an interaction in which villain status reduces discrimination between weak and strong case-specific evidence, reduces willingness to revise, or increases vice attribution toward a principled intervener.

The clearest contrary result would be stable sensitivity to evidence, stable openness to correction, and no added hostility toward the rational intervener after accounting for genuine obstruction. The framework is useful only if the experiment gives it a fair chance to fail.

C.6 Participant protection

Because the study concerns accusation, stigma, and punishment, researchers should avoid real targets and graphic allegations when a less harmful vignette can test the same mechanism. Debriefing should state that the people and events were fictional, explain why evidence strength and social feedback were manipulated, and provide a way to withdraw data where ethics rules require it.